TINA: Mathematical Companion

Scott Moeller Bhaskar Krishnamachari

September 2026

Mathematical companion to *A Theory of Information Architecture for Networked Decisions: Freshness, Locality, and Coordination*, the alternative accessibility revision of arXiv v2. Unqualified references to main-paper results and equations use the accompanying paper's numbering.

1 Mathematical Toolkit

This optional companion collects background for the paper and its technical appendices. The main composition proof needs conditional expectation and error expansion, not Gaussianity or projection theory. All random vectors in the projection statements are square-integrable, and all weighting matrices Q are deterministic and positive definite.

1.1 Information as σ-algebras

σ-algebra generated by a variable. For a random vector Y, $\sigma(Y)$ is the collection of all events whose occurrence can be decided by looking at Y. It formalizes "what is known if one observes Y."

Measurability means "a function of". By the Doob–Dynkin lemma, a random variable u is $\sigma(Y)$-measurable if and only if $u = g(Y)$ for some (measurable) function g. So "u_i is $\mathcal{G}_i$-measurable" means "u_i is computed from agent i's observations."

Combining information. $\sigma(Y) \vee \sigma(Z) = \sigma(Y, Z)$ is the information from observing both.

Nesting. $\mathcal{G} \subseteq \mathcal{G}'$ means that whoever knows $\mathcal{G}'$ knows at least everything in $\mathcal{G}$. Every $\mathcal{G}$-measurable rule is then also $\mathcal{G}'$-measurable.

1.2 Geometry of square-integrable random vectors

The space L^2. Random vectors u with $\mathbb{E}\|u\|^2 < \infty$, with inner product $\langle u, v\rangle_Q = \mathbb{E}[u^\top Q v]$ for a fixed positive-definite Q. This is a Hilbert space: an inner-product space in which limits of Cauchy sequences stay in the space.

Projection theorem. For a closed linear subspace S and any x, there is a unique $\Pi x \in S$ closest to x. It is characterized by orthogonality: $\langle x - \Pi x,\, s\rangle_Q = 0$ for all $s \in S$.

Pythagoras. $\|x\|^2 = \|\Pi x\|^2 + \|x - \Pi x\|^2$. More generally, for any $s \in S$, $\|x - s\|^2 = \|x - \Pi x\|^2 + \|\Pi x - s\|^2$.

Conditional expectation is a projection. $\mathbb{E}[x \mid \mathcal{G}]$ is the orthogonal projection of x onto $\{\mathcal{G}\text{-measurable } L^2 \text{ variables}\}$. It is therefore the best mean-square predictor of x among *all* functions of the information, not only linear ones.

Tower property. For square-integrable $g(Y)$ and Z, $\mathbb{E}[g(Y)^\top Z] = \mathbb{E}\big[g(Y)^\top \mathbb{E}[Z \mid Y]\big]$. *Consequence:* if $\mathbb{E}[\varepsilon \mid Y] = 0$, then ε is orthogonal to every function of Y. This is the step used in the proof of Theorem 1.

Common information makes the optimizer independent of Q. If every component of the policy may depend on the same σ-algebra $\mathcal{G}$, the Q-projection of x is simply $\mathbb{E}[x \mid \mathcal{G}]$, for any positive-definite Q. The error $e = x - \mathbb{E}[x \mid \mathcal{G}]$ satisfies $\mathbb{E}[s^\top Q e] = \mathbb{E}\big[s^\top Q\, \mathbb{E}[e \mid \mathcal{G}]\big] = 0$ for every $\mathcal{G}$-measurable s. With *different* information per agent, this fails and the coupled equations (A.14) are needed.

Regret and the factor one-half. Completing the square in the main paper's quadratic cost gives $R(\mathcal{A}) = \frac{1}{2}\|x^\star - u^\star_{\mathcal{A}}\|^2_Q$. Consequently a squared Q-norm in a proof equals twice the corresponding regret. The best constant action is $\bar{x}^\star = \mathbb{E}[x^\star]$, since

$$\|x^\star - c\|^2_Q = \|x^\star - \bar{x}^\star\|^2_Q + (\bar{x}^\star - c)^\top Q(\bar{x}^\star - c).$$

Thus $2R_\infty = \|x^\star - \bar{x}^\star\|^2_Q$.

Mean-square convergence. $Y_n \to Y$ in L^2 means $\mathbb{E}\|Y_n - Y\|^2 \to 0$.

1.3 Gaussian facts

Conditioning. If (X, Y) is jointly Gaussian and Σ_{YY} is invertible,

$$\mathbb{E}[X \mid Y] = \mu_X + \Sigma_{XY}\Sigma_{YY}^{-1}(Y - \mu_Y), \qquad \mathrm{Cov}(X \mid Y) = \Sigma_{XX} - \Sigma_{XY}\Sigma_{YY}^{-1}\Sigma_{YX}.$$

The conditional mean is affine in Y (linear after centering). For singular Σ_{YY}, the same formulas use its Moore–Penrose inverse on the covariance support. The conditional covariance (a *Schur complement*) does **not** depend on the observed value of Y.

Scalar case. For nonzero marginal variances, $\mathrm{Var}(X \mid Y) = \mathrm{Var}(X)\big(1 - \mathrm{Corr}(X, Y)^2\big)$.

Uncorrelated implies independent for jointly Gaussian variables. For the OU process, Brownian increments after the snapshot time are independent of the past, which gives innovation independence in Section 1.4.

Linear functionals stay Gaussian. Integrals such as $\int w(u)\theta(u)\,du$ defined as mean-square limits of linear combinations of a Gaussian field are Gaussian, so the facts above apply to $x^\star$ in Section 4.3.

1.4 The Ornstein–Uhlenbeck process in time

Scalar SDE. $d\theta_t = -\frac{1}{T}(\theta_t - \bar{\theta})\,dt + \sigma\sqrt{2/T}\,dW_t$, where W is Brownian motion (continuous paths, independent Gaussian increments with variance equal to elapsed time). A useful mental model is the small-step recursion

$$\theta_{t+\Delta} \approx \theta_t - \tfrac{\Delta}{T}(\theta_t - \bar{\theta}) + \sigma\sqrt{2\Delta/T}\,Z, \qquad Z \sim N(0,1) \text{ fresh at each step.}$$

Solution over a lag τ.

$$\theta_t - \bar{\theta} = e^{-\tau/T}(\theta_{t-\tau} - \bar{\theta}) + \underbrace{\sigma\sqrt{2/T}\int_{t-\tau}^{t} e^{-(t-s)/T} dW_s}_{\varepsilon_{t,\tau}}.$$

The innovation is Gaussian with mean zero, independent of the past, and has variance $\sigma^2\frac{2}{T}\int_0^\tau e^{-2s/T}ds = \sigma^2(1 - e^{-2\tau/T})$. This shows why the factor $\sqrt{2/T}$ makes the stationary variance equal σ^2.

Correlation. In stationarity, $\mathrm{Corr}(\theta_t, \theta_s) = e^{-|t-s|/T}$. The coherence time T is the lag at which correlation falls to $1/e \approx 0.37$.

Multivariate, common rate (A.15)**.** Replace σ by $\Sigma_S^{1/2}$, the symmetric positive-semidefinite matrix whose square is Σ_S. Every formula above holds with σ^2 replaced by Σ_S, because the drift $-\frac{1}{T}I$ treats all directions alike.

Markov property. The future after time s depends on the past only through θ_s.

Discrete-time analogue. AR(1): $\theta_k - \bar{\theta} = a(\theta_{k-1} - \bar{\theta}) + \varepsilon_k$. In its stationary solution with $|a| < 1$ and i.i.d. square-integrable innovations independent of the past, each nondegenerate scalar component has correlation $a^{|h|}$ at lag h.

1.5 Exponential covariance in space; fields on graphs

1-D spatial OU. A Gaussian field on the line with covariance $\sigma^2 e^{-|z-z'|/\ell_s}$ is an OU process in the spatial coordinate. It is Markov in z: given $\theta(z_0)$, the field on either side of z_0 is conditionally independent of the other side. Given a whole segment, the outside depends only on the nearer endpoint.

Gaussian Markov random fields (GMRFs). On an undirected graph, a Gaussian vector with nonsingular covariance is Markov with respect to the graph when each node, given its neighbors, is independent of all other nodes. Equivalently, the precision matrix Σ^{-1} has zeros for non-adjacent pairs.

Graph Laplacian. $L_G = D_G - A_G$, where D_G is the diagonal degree matrix and A_G the adjacency matrix of a simple undirected unweighted graph. For a signal x, $x^\top L_G x = \sum_{(i,j)\in E}(x_i - x_j)^2$ measures roughness. Eigenvectors of L_G with small eigenvalues ("low graph frequencies") vary slowly across edges.

Covariance (39)**.** $\Sigma_S = \sigma^2(\kappa_s^2 I + L_G)^{-\nu}$ puts large variance on low-frequency (smooth) patterns. After normalizing total variance, smaller κ_s or larger ν shifts relative energy toward smoother modes. This spectral model need not have a sparse precision matrix for arbitrary ν, so it need not be Markov on the original graph.

1.6 Matrix facts

Quadratic-form expectation. For a random vector Y with mean μ, $\mathbb{E}[Y^\top A Y] = \operatorname{tr}(A \operatorname{Cov} Y) + \mu^\top A \mu$.

Cyclic trace. $\operatorname{tr}(ABC) = \operatorname{tr}(BCA) = \operatorname{tr}(CAB)$ whenever the products are defined.

Sherman–Morrison. For invertible A and vectors a, c with $1 + c^\top A^{-1} a \neq 0$,

$$(A + ac^\top)^{-1} = A^{-1} - \frac{A^{-1} a c^\top A^{-1}}{1 + c^\top A^{-1} a}.$$

With $A = qI$ and $a = \gamma\mathbf{1}$, $c = \mathbf{1}$ this gives $(qI + \gamma\mathbf{1}\mathbf{1}^\top)^{-1} = \frac{1}{q}\big(I - \frac{\gamma}{q+\gamma n}\mathbf{1}\mathbf{1}^\top\big)$. Also $\operatorname{tr}(qI + \gamma\mathbf{1}\mathbf{1}^\top) = n(q+\gamma)$.

Positive-definite weighting. $Q \succ 0$ means $x^\top Q x > 0$ for all $x \neq 0$. Then $\lambda_{\min}(Q)\|x\|^2 \leq x^\top Q x \leq \lambda_{\max}(Q)\|x\|^2$, so the Q-norm and the ordinary norm are equivalent.

Null space and modes. The null space of K is $\{\xi : K\xi = 0\}$. A "mode" of the state is a direction ξ, often an eigenvector of Σ_S. Its decision-error weight per unit state variance is $(K\xi)^\top Q (K\xi)$.

1.7 Calculus, asymptotics, and optimization

Log-derivative. $\frac{d}{dr}\log f(r) = f'(r)/f(r)$ is the *relative* growth rate of f, the fractional change per unit r. The marginal-balance condition compares the log-derivative of the fresh captured share

with the negative log-derivative of the temporal survival factor.

First-order conditions. At an interior minimum of a differentiable function the derivative vanishes. The converse fails: a zero derivative identifies candidates only.

Single crossing. If $\Delta(r)$ is continuous and strictly decreasing, it has at most one zero. The function whose derivative has the sign of $-\Delta$ is then unimodal, with its minimum at that zero, or at an endpoint if there is no zero.

Continuous relaxation. Solving an integer problem (radius in hops) by first allowing real values. When the continuous objective is unimodal, the integer optimum is the better of the two integers bracketing the continuous optimum.

Taylor approximations used. $1 - e^{-x} \approx x$ and $\log(1 + x) \approx x$ for small x.

Asymptotic notation. $f = O(g)$ means $|f| \leq C|g|$ near the limit; $f = o(g)$ means $f/g \to 0$.

Positive part and extended reals. $[z]_+ = \max\{z, 0\}$. Writing $\rho^\star = +\infty$ means that no finite value satisfies the condition.

2 Glossary

The entries below use the meanings and assumptions of this paper.

Term	Meaning in this paper	First use/background
Information architecture	Specification of which variables, at which ages, each agent's decision may depend on; formally a tuple of σ-algebras $(\mathcal{G}_1, \ldots, \mathcal{G}_n)$	Section 2.2
σ-algebra $\sigma(Y)$	"What is known from observing Y"	Appendix A, Section 1.1
Measurable ($\mathcal{G}_i$-measurable)	Computable from agent i's information	(A.6), Section 1.1
Admissible policy	A joint rule in which each agent uses only its own information	(A.6)
Static team	Several decision makers with one shared objective, each acting on its own information, whose actions do not affect anyone's later information	Section 1, Section 6.1
Person-by-person (stationarity) conditions	Each agent's rule is optimal given the others' rules; here, equations (A.13)–(A.14)	Corollary 2
Full-information (clairvoyant) decision	$x_t^\star = Q^{-1}b(\theta_t)$, optimal with complete current state	(A.3)
Architecture regret	Extra expected cost versus the full-information decision, after optimizing over admissible policies	(5)
Q-weighted norm	$\|u\|_Q^2 = \mathbb{E}[u^\top Q u]$; the metric in which regret is a squared distance	(A.9)
Projection	Closest point in a subspace; here, the best admissible policy	Lemma 2, Section 1.2
No-information	Using no state information; the best constant action	(6)
No-information regret R_∞	Regret of the best constant; the normalizer for all shares	(6)
Omission share η_S	Fraction of decision variance not captured by current information in a given scope	(7)
Staleness ratio ρ	τ/T, delay measured in coherence times	Section 2.4
Stationary	Joint distributions are invariant under time (or spatial) translations	(A.1)
Ergodic	Time averages along one realization converge to expectations	Appendix A
Square-integrable	Finite second moment, $\mathbb{E}\|Y\|^2 < \infty$	(A.1)
Gauss–Markov	Gaussian and Markov; the OU process is a sufficient special case	Appendix A.1
Ornstein–Uhlenbeck (OU) process	Stationary Gaussian Markov process that reverts to its mean with exponential correlation	(A.15), Section 1.4
Coherence time T	Time lag at which correlation falls to $1/e$	(A.15)
Common-rate model	Conditional mean of the whole current state given the old state equals $e^{-\tau/T}$ times the old state	(8)
Innovation	Part of the current state not predictable from an earlier snapshot	(8)
Mixing	Decay of statistical dependence with separation; not a standing assumption of the main theorem	Section 2.4

Term	Meaning in this paper	First use/background
Data-processing inequality	Processing information along a Markov chain cannot improve prediction	Section 1.4
Nested (architectures)	One architecture's information contains the other's, at every agent	(A.7), Section 3
Synchronized snapshot	All observations used in a decision share one time stamp	Section 2.2
Mode	A direction or spatial pattern in state space, often an eigenvector of Σ_S	Section 2.1
Null space of K	State directions that do not change the desired decision	Section 2.1, Section 1.6
Decision-sensitivity operator K	$Q^{-1}B$; maps state to desired action	(A.5)
Separable covariance	Space–time covariance that factors into a time part times a space part	(A.16)
Spatial correlation length ℓ_s	Distance over which the state stays correlated	(31)
Decision-relevance length ℓ_c	Distance over which remote state affects the desired local action	(32)
Temporal propagation length L_T	vT: how far information travels in one coherence time	(35)
Spatial Markov property	Given the field at a point, the two sides are conditionally independent	Section 4.3, Section 1.5
Gaussian Markov random field	Gaussian vector with graph-based conditional independence	Section 1.5
Graph Laplacian L_G	$D_G - A_G$; measures roughness of signals on a graph	(39), Section 1.5
Graph frequency	Eigenvalue of L_G; low values correspond to smooth patterns	Section 6.5, Section 1.5
Continuous relaxation	Treating the integer radius as a real number	Section 4
Unimodal	Decreasing, then increasing; a single minimum	Section 4
Marginal balance	Equality of spatial-gain and freshness-loss rates, necessary at a differentiable interior optimum	(29)
Comparative statics	How the optimum moves when a parameter changes	Theorem 3
Extended reals	Real numbers together with $\pm\infty$	Theorem 2
Age of Information (AoI)	Time since the freshest received update was generated	Section 6.3
Witsenhausen's counterexample	Classic example where control actions also signal information, and linear policies can be suboptimal	Section 1, Section 7

3 Snapshot timing

The diagram illustrates the main paper's snapshot convention. Section 2 of the main paper states the non-Gaussian temporal assumptions; Section 4.3 adds the Gaussian spatial specialization.

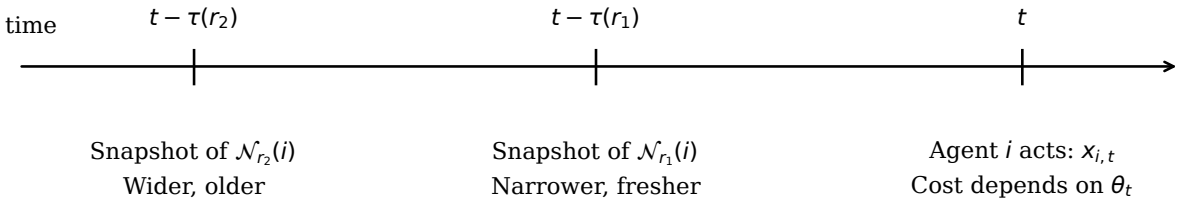


Figure 1: A radius-r synchronized snapshot is taken at time $t - \tau(r)$ and used for the decision at time t. For $r_2 > r_1$ with $\tau(r_2) > \tau(r_1)$, the wider snapshot is older. Cost is evaluated against the current state θ_t.

4 Marginal-rate illustration

Figure 2 illustrates the marginal-balance condition for the canonical model. In this diagram only, write $m_S = (d/dr)\log[1-\eta_S(r)]$ and $m_T = 2\tau'(r)/T$.

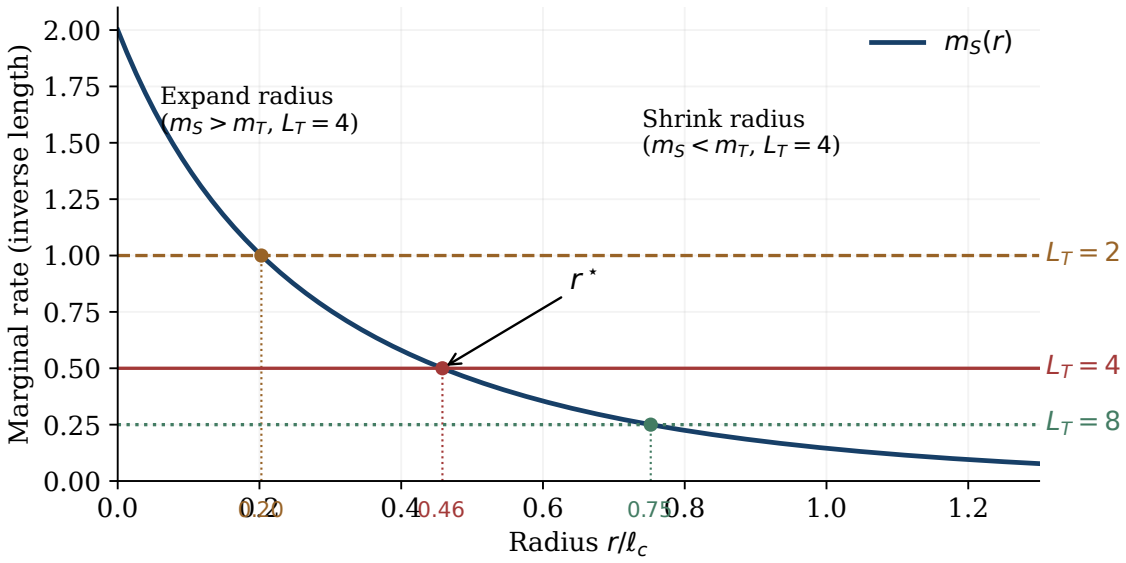


Figure 2: Marginal balance in the canonical model with $\ell_c = 1$ and $\ell_s = 0.5$. Increasing $L_T = vT$ lowers $m_T = 2/L_T$ and moves the crossing with m_S outward. For the middle curve ($L_T = 4$), regret falls to the left of the crossing and rises to its right. This is an illustration of the analytical rates, not an additional numerical experiment.