A Theory of Information Architecture for Networked Decisions: Freshness, Locality, and Coordination

Scott Moeller
Independent Researcher, Portsmouth, RI, USA
scott.h.moeller@gmail.com

Bhaskar Krishnamachari
University of Southern California
bkrishna@usc.edu

September 2026

Abstract

A broader view of a network can help agents coordinate, but collecting that view takes time. When is a fresh local observation more useful than an older view of the whole network, and when should an agent look only partway across it? We study agents that repeatedly choose actions to minimize a shared quadratic cost in an environment that evolves independently of their actions. Each agent observes a snapshot of its neighborhood taken at a common time. When the full-information action is affine in the state and the conditional mean of the current state is the old state scaled by a common exponential decay factor, we obtain an exact rule: the reduction in cost provided by a delayed snapshot equals the reduction provided by a fresh snapshot of the same neighborhood, multiplied by the square of the decay factor. The result allows non-Gaussian states and nonlinear decision rules. It gives the delay at which fresh local information overtakes a stale global view. In a Gaussian spatial model, it also gives the optimal observation radius in terms of how far state correlations extend, how far away observations matter to the decision, and how far information travels in one coherence time. On a fixed 64-node ring, changing only the decision being made shifts the optimal radius from one to five hops.

1 Introduction

Many networked systems must repeatedly make coupled decisions using information that is distributed across space and changes over time. A wireless transmitter may know its current local channel state but not the current interference conditions elsewhere in the network; an edge server may observe its own workload immediately while a system-wide view is available only after telemetry has been collected and disseminated; and a distributed sensing system may have fresh measurements from nearby sensors while observations from distant sensors arrive with greater latency. In each case, broad information can improve coordination, but obtaining it takes time. The resulting design choice is not simply between centralized and distributed control, but between *information architectures* that differ simultaneously in spatial scope and temporal freshness.

A two-sensor example. Two sensors report independent signals X_t and Y_t, each with mean zero and variance σ^2. We want to estimate their current average $Z_t = (X_t + Y_t)/2$. Which is more useful: one sensor's reading now, or both sensors' readings from τ seconds ago?

With only X_t, our best estimate is $X_t/2$. The error is the missing term $Y_t/2$, so the mean squared error is $\sigma^2/4$. With both old readings, we know the old average $Z_{t-\tau}$ exactly. Suppose the two signal processes are independent and stationary, and each satisfies the scalar conditional-mean

prediction rule $\mathbb{E}[X_t \mid X_{t-\tau}] = e^{-\tau/T} X_{t-\tau}$ (and likewise for Y), where T is the coherence time. Our best estimate today is then $e^{-\tau/T} Z_{t-\tau}$. Its mean squared error is $\mathrm{Var}(Z_t)(1 - e^{-2\tau/T})$. Since $\mathrm{Var}(Z_t) = \sigma^2/2$, the fresh reading wins when

$$\frac{\sigma^2}{4} < \frac{\sigma^2}{2}(1 - e^{-2\tau/T}) \quad \Longleftrightarrow \quad \tau > \frac{T \log 2}{2} \approx 0.35T.$$

One fresh sensor leaves half the variation in the desired average unexplained. Once the old pair leaves more than half unexplained, the fresh reading is more useful. Theorem 2 gives the corresponding crossover when several agents choose actions that interact through a shared cost.

This tension appears in several literatures under different performance models. Work on stale information in distributed systems shows that old global state can lose much of its value [1], and Age of Information formalizes freshness as a system resource [2]. Distributed optimization studies computation under communication constraints [3], organizational economics studies adaptation to local information versus coordination [4], and networked-control work examines how communication scope, topology, and propagation delay interact with closed-loop performance.

Distributed-control work establishes controller localization and models finite-speed or distance-dependent communication [5–8]. Most directly, Ballotta, Jovanović, and Schenato show that scope-dependent delay can make sparse feedback outperform all-to-all control [9]; Section 6.2 develops the comparison with this literature.

We consider agents whose actions incur a shared quadratic cost in an environment that evolves independently of those actions. An information architecture specifies which state observations each agent can use and how old they are. We compare a fresh local reading, a delayed global view, and synchronized neighborhood snapshots by how well they support the current full-information decision. Current actions affect neither future state nor any agent's later information; this is a repeated static-team problem.

We build on the classical team theory of Radner and of Marschak and Radner [10, 11], which asks how agents with different information can best pursue a shared objective. We use the best decision rules available to the team under each information architecture to quantify how much is gained by observing a wider neighborhood and how much is lost when those observations arrive later. For the main results, each agent acts on a synchronized snapshot: widening scope can replace a fresh local reading by an older one. Architectures that retain fresh local data while adding remote data have different information and fall outside this comparison.

In the shared-resource example, stronger decision coupling makes a global view worth waiting longer for. The choice of scope also depends on which state changes matter to the decision, not only on how strongly nearby states are correlated.

The contributions are:

1. For affine desired actions under a common scalar conditional-mean prediction rule, we prove a spatial–temporal composition law for all admissible rules and coupled quadratic costs. Its endpoint consequences give the fresh-local versus stale-global crossover.

2. We characterize interior radius optima by marginal balance and give a global solution for exponential spatial omission. A derived omission curve in a Gaussian spatial model connects this solution to three physical lengths: spatial correlation, decision relevance, and temporal propagation.

3. We illustrate how decision-relevant predictability governs architecture choice on fixed and noncanonical finite graphs. The coupling calculation and canonical spatial derivation supply information beyond the composition law itself.

Section 2 defines observations, regret, and the temporal prediction assumption. Section 3 proves the composition law and derives the endpoint comparison and coupling example. Section 4 chooses the radius. Section 5 gives four numerical illustrations, including noncanonical graphs; Sections 6 and 7 discuss related work and implications. Technical appendices retain the broader formulation, proofs, and reproducibility details.

2 Model and observations

We compare what agents can achieve with different observations of the same environment. The environment evolves independently of their actions, and actions do not affect later observations. Each epoch is therefore a static team decision. Actions are unconstrained Euclidean vectors.

2.1 The desired action and its cost

There are n agents on a graph $G = (V, E)$. Its hop distance $d_G(i, j)$ defines the neighborhood

$$\mathcal{N}_r(i) = \{j \in V : d_G(i, j) \le r\}. \tag{1}$$

The graph describes observation scope and acquisition delay. The stacked state $\theta_t = (\theta_{1,t}, \dots, \theta_{n,t})$ has finite second moment; local states and local actions may each be vectors. Write $\Sigma_S = \mathrm{Cov}(\theta_t)$ for the stationary state covariance. We center known means throughout the main argument. Subtracting the mean state and mean desired action is a change of coordinates, not a restriction to physically zero-mean environments.

With complete current information, the desired action is $x_t^\star = K\theta_t$. We measure excess cost by the coupled tracking loss

$$x_t^\star = K\theta_t, \qquad \ell(x, \theta_t) = \frac{1}{2}(x - x_t^\star)^\top Q(x - x_t^\star), \qquad Q \succ 0. \tag{2}$$

Here K is a fixed matrix describing how state affects the desired action, and Q is deterministic, symmetric, and positive definite. Its off-diagonal blocks allow agents' action errors to interact in the shared cost. Neither matrix must follow the communication graph. Completing the square in a quadratic objective with affine state dependence gives exactly this loss after centering; Appendix A gives that equivalence and the broader fixed-epoch formulation with non-affine desired actions.

Decision relevance. A state change ξ changes the desired action by $K\xi$. If $K\xi = 0$, that change creates no decision error. Other directions matter according to $(K\xi)^\top Q(K\xi)$. In particular,

$$\mathrm{Cov}(x_t^\star) = K\Sigma_S K^\top. \tag{3}$$

The value of observations depends jointly on K, Q, the state covariance, and which states are observed. Figure 3 illustrates this dependence with the graph and state covariance held fixed.

2.2 What each agent observes

At radius r and age τ, agent i observes the vector

$$Z_{i,t}^{(r,\tau)} = (\theta_{j,t-\tau} : j \in \mathcal{N}_r(i)). \tag{4}$$

We denote this architecture by $\mathcal{A}_{r,\tau}$. The two parameters are independent here; Section 4 later sets $\tau = \tau(r)$. The three cases are:

- **Fresh local:** each agent observes its current state $\theta_{i,t}$, giving $\mathcal{A}_{0,0}$ and regret R_{loc}.
- **Stale global:** every agent observes the same complete old state $\theta_{t-\tau}$, giving regret $R_{\mathrm{glob}}(\tau)$.
- **Synchronized neighborhood:** each agent observes $Z_{i,t}^{(r,\tau)}$, so every reading used has the same timestamp.

An admissible rule is $u_i = g_i(Z_{i,t}^{(r,\tau)})$ with finite second moment. We allow all such rules, including nonlinear ones.

Widening a synchronized snapshot may replace nearby fresh readings with older readings so that all observations have a common timestamp. Retaining fresh local readings while adding old remote readings gives a different architecture. Adding usable information without discarding existing information cannot increase optimal cost, because agents can ignore it. The same observation applies to a delayed global summary added to local data; such mixed-age observations are outside the snapshot family studied here.

2.3 Regret and the fresh omission share

Regret is the smallest expected excess cost achievable when each agent uses only its specified observations:

$$R(\mathcal{A}) = \inf_{\{g_i\}} \frac{1}{2}\mathbb{E}[(u - x_t^\star)^\top Q(u - x_t^\star)]. \tag{5}$$

This optimizes over admissible rules, not a particular distributed algorithm. It uses the fixed-information static-team benchmark [10, 12]; the formal existence, uniqueness, and projection result is in Appendix A. With a common observation, the optimum is the vector conditional mean of $x_t^\star$. With different observations and off-diagonal Q, separate agentwise conditional means generally do not minimize team cost.

The best no-information action is zero in centered coordinates. Its loss is

$$R_\infty = \frac{1}{2}\mathbb{E}[(x_t^\star)^\top Q x_t^\star] = \frac{1}{2}\operatorname{tr}(K^\top Q K \Sigma_S). \tag{6}$$

For $R_\infty > 0$, define

$$\eta_S(r) = \frac{R(\mathcal{A}_{r,0})}{R_\infty}, \qquad 0 \le \eta_S(r) \le 1. \tag{7}$$

This is the fraction of no-information decision loss remaining with a fresh radius-r view. Larger fresh neighborhoods cannot increase it, since their extra observations can be ignored. Radius zero gives $\eta_S(0) = R_{\mathrm{loc}}/R_\infty$; at the diameter of a finite connected graph the fresh view is complete and omission is zero.

2.4 A common temporal prediction rule

Assume a stationary state process with finite second moments and, for every delay analyzed, the transition

$$\theta_t = a\theta_{t-\tau} + \varepsilon_{t,\tau}, \qquad a = e^{-\tau/T}, \quad T > 0, \qquad \mathbb{E}[\varepsilon_{t,\tau} \mid \theta_{t-\tau}] = 0. \tag{8}$$

The best mean prediction of current state from the complete old state is the old state multiplied by a common decay factor. The remaining disturbance has mean zero even after the *whole* old state is known. This assumption requires neither Gaussian states nor independence of the disturbance from the old state. State coordinates may be spatially dependent. We write $\rho = \tau/T$ for delay measured in coherence times.

Stationarity supplies the same one-time distribution at the two epochs. Equality of covariance matrices alone would not identify optimal errors over arbitrary nonlinear rules. Likewise, a covariance decay relation alone does not imply the conditional-mean property in Eq. (8).

Imagine two edge servers whose average workload changes slowly while their workload difference changes quickly. An old report can remain useful for total demand after losing value for assigning the next job. Our common-rate assumption assigns both patterns the same decay time T; distinct rates are outside the exact factorization below. A stationary Gaussian Ornstein–Uhlenbeck process is one sufficient special case [13], described in Appendix A.1. A stationary AR(1) process with coefficient $0 < c < 1$ and independent, mean-zero finite-variance innovations gives the analogous prediction coefficient c^k at integer lag k.

3 How scope and freshness combine

The temporal prediction rule separates the loss from new disturbances from the loss due to incomplete spatial information. This separation applies to the joint team optimization even when actions are coupled.

Theorem 1 (Spatio-temporal composition)**.** *Under the centered affine tracking model* (2)*, stationary conditional-mean transition* (8)*, and synchronized observations* (4)*, every radius r and finite delay $\tau \geq 0$ satisfy*

$$\boxed{R(\mathcal{A}_{r,\tau}) = (1-a^2)R_\infty + a^2 R(\mathcal{A}_{r,0}), \qquad a = e^{-\tau/T}.} \tag{9}$$

The optimization allows all admissible finite-second-moment rules and any deterministic $Q \succ 0$. For $R_\infty > 0$, equivalently,

$$\boxed{1 - \frac{R(\mathcal{A}_{r,\tau})}{R_\infty} = e^{-2\tau/T}[1 - \eta_S(r)].} \tag{10}$$

The fraction of no-information loss eliminated equals the fraction captured by a fresh neighborhood times the fraction of predictive value surviving the delay. Spatial covariance is unrestricted under the stated temporal assumption.

Proof. Applying K to the transition gives

$$x_t^\star = a x_{t-\tau}^\star + K\varepsilon_{t,\tau}. \tag{11}$$

Every joint admissible rule u is a function of the complete old state, although its components each use only their own neighborhood. On expanding the squared tracking error, the cross term vanishes:

$$\begin{aligned}&\mathbb{E}[(a x_{t-\tau}^\star - u)^\top Q K \varepsilon_{t,\tau}] \\ &\quad = \mathbb{E}[(a x_{t-\tau}^\star - u)^\top Q K \mathbb{E}[\varepsilon_{t,\tau} \mid \theta_{t-\tau}]] = 0.\end{aligned}$$

Thus the error is the sum of a predictable-part error and a disturbance error:

$$\begin{aligned}\mathbb{E}[(x_t^\star - u)^\top Q (x_t^\star - u)] = \mathbb{E}[(a x_{t-\tau}^\star - u)^\top Q (a x_{t-\tau}^\star - u)] \\ + \mathbb{E}[(K\varepsilon_{t,\tau})^\top Q (K\varepsilon_{t,\tau})].\end{aligned} \tag{12}$$

Apply the same expansion with $u = 0$. Stationarity gives

$$\mathbb{E}[(K\varepsilon_{t,\tau})^\top Q (K\varepsilon_{t,\tau})] = (1-a^2)\mathbb{E}[(x_t^\star)^\top Q x_t^\star] = 2(1-a^2)R_\infty. \tag{13}$$

For finite delay $a > 0$, write $u = a\widetilde{u}$. This rescaling is a bijection of the admissible rule class: each component uses exactly the same observations and still has finite second moment. Consequently,

$$\begin{aligned}&\inf_{u} \mathbb{E}[(ax^\star_{t-\tau} - u)^\top Q(ax^\star_{t-\tau} - u)] \\ &\quad = a^2 \inf_{\widetilde{u}} \mathbb{E}[(x^\star_{t-\tau} - \widetilde{u})^\top Q(x^\star_{t-\tau} - \widetilde{u})] = 2a^2 R(\mathcal{A}_{r,0}). \end{aligned} \tag{14}$$

The last equality uses stationarity and the same observation pattern at each epoch. Combining the two errors and dividing by two proves the law; normalization is valid when $R_\infty > 0$. □

The proof rescales a joint rule; it does not replace a coupled team optimum by separate conditional means. The disturbance since the snapshot is unpredictable from any of its neighborhoods. A wider neighborhood can reduce only the remaining, temporally predictable part of the error.

3.1 Complete old information versus current local information

Corollary 1 (Exact stale-global regret)**.** *Under Theorem 1, a complete snapshot of age τ gives*

$$R_{\mathrm{glob}}(\tau) = (1 - e^{-2\tau/T})R_\infty. \tag{15}$$

Proof. With a complete fresh snapshot the desired action is feasible, so $R(\mathcal{A}_{r,0}) = 0$ in Eq. (9). □

The optimal global action is $aK\theta_{t-\tau}$. Its regret is zero at zero delay, grows initially as $2(\tau/T)R_\infty + o(\tau/T)$, and tends to R_∞ as delay grows. This last statement is a limit, not an invertible policy rescaling at infinite delay.

If different parts of the state change at different rates, a complete old snapshot can still have an explicit value. Suppose it reveals independent stationary Ornstein–Uhlenbeck modes with coherence times T_k, and let w_k be the fraction of the no-information decision loss $R_\infty > 0$ attributable to mode k, with $\sum_k w_k = 1$. Then

$$1 - \frac{R_{\mathrm{glob}}(\tau)}{R_\infty} = \sum_k w_k e^{-2\tau/T_k}.$$

In the two-server example in Section 2.4, an old snapshot remains useful longer for estimating total workload if total workload changes more slowly than the difference between the two loads. This formula applies to a complete old snapshot. Extending it to partial neighborhood observations requires additional assumptions.

Theorem 2 (Fresh-local versus stale-global crossover)**.** *Under Theorem 1, let $R_\infty > 0$ and $0 \le \eta_S(0) < 1$. Fresh-local information is strictly better exactly when*

$$\boxed{\frac{\tau}{T} > \rho^\star \triangleq \frac{1}{2}\log\frac{1}{1-\eta_S(0)}.} \tag{16}$$

Equality gives equal regret. If $\eta_S(0) = 1$, no finite delay gives a strict local advantage; we write $\rho^\star = +\infty$.

Proof. Divide the two regrets by R_∞: the comparison is $\eta_S(0) < 1 - e^{-2\tau/T}$. Rearranging gives the threshold and equality case. □

When $\eta_S(0) = 0$, the local view already determines the desired action, and every positive global delay makes local information preferable. For an omission share strictly between zero and one, the two losses cross once. The threshold concerns prediction of the desired action, rather than raw state correlation alone.

3.2 Role of Decision Coupling

The fresh-local omission share $\eta_S(0)$ depends jointly on the spatial statistics of the environment and on the coupling structure of the decision problem. To isolate the role of decision coupling, consider a symmetric shared-resource model with $n \geq 2$ scalar actions,

$$f(x,\theta) = \sum_{i=1}^{n} \left(\frac{q}{2} x_i^2 - \theta_i x_i \right) + \frac{\gamma}{2} \left(\sum_{i=1}^{n} x_i \right)^2, \qquad q > 0,\ \gamma \geq 0. \tag{17}$$

Here q determines the local curvature or stiffness of each decision, while γ determines the strength of the shared coordination penalty. The cost has

$$Q = qI + \gamma \mathbf{1}\mathbf{1}^\top, \qquad K = Q^{-1}. \tag{18}$$

Here $\mathbf{1}$ is the all-ones vector.

Assume that the state components θ_i are independent and mean-zero, with equal finite variance $\sigma^2 > 0$. This calculation concerns a single epoch.

The full-information action $x^\star = Q^{-1}\theta$ has components

$$x_i^\star = \frac{\theta_i}{q} - \frac{\gamma}{q(q+\gamma n)} \sum_{j=1}^{n} \theta_j. \tag{19}$$

Thus, even though the linear reward $\theta_i x_i$ is local, coupling through the common resource makes agent i's optimal action depend on the states of all other agents.

Proposition 1 (Local explainability under aggregate coupling)**.** *For the model* (17)*, the optimal fresh-local policy is*

$$u_i^{\mathrm{loc}} = \frac{\theta_i}{q+\gamma}, \tag{20}$$

and the fresh-local omission share is

$$\boxed{\eta_S(0) = \frac{\gamma^2(n-1)}{(q+\gamma)\,(q+\gamma(n-1))}.} \tag{21}$$

Proof. For any admissible $u_i = g_i(\theta_i)$, write $u_i = \mu_i + \widetilde{u}_i$, where $\mu_i = \mathbb{E}[u_i]$ and $\mathbb{E}[\widetilde{u}_i] = 0$. Independence across agents and zero state means give

$$\mathbb{E}[f(u,\theta)] = \sum_i \mathbb{E}\left[\frac{q+\gamma}{2} \widetilde{u}_i^2 - \theta_i \widetilde{u}_i \right] + \frac{q}{2} \sum_i \mu_i^2 + \frac{\gamma}{2} \left(\sum_i \mu_i \right)^2. \tag{22}$$

The mean terms are minimized at zero. Each remaining term is minimized pointwise by $\widetilde{u}_i = \theta_i/(q+\gamma)$, which is feasible and has zero mean. Thus the expected local cost is $-n\sigma^2/[2(q+\gamma)]$, and the best constant cost is zero. The eigenvalues of Q are $q+\gamma n$ once and q with multiplicity $n-1$. The full-information expected cost is therefore $-(\sigma^2/2)[(n-1)/q + 1/(q+\gamma n)]$. Subtracting this cost from the local and constant costs gives Eqs. (B.1)–(B.2) in Appendix B; their ratio is the stated omission share. □

Under the common-rate temporal assumption, combining Proposition 1 with Theorem 2 gives an explicit threshold.

$$\boxed{\rho^\star(\gamma, n) = \frac{1}{2}\log\left[\frac{(q+\gamma)\,(q+\gamma(n-1))}{q(q+\gamma n)}\right].} \tag{23}$$

Therefore

$$R_{\mathrm{loc}} < R_{\mathrm{glob}}(\tau) \quad \Longleftrightarrow \quad \frac{\tau}{T} > \rho^\star(\gamma, n). \tag{24}$$

As $n \to \infty$,

$$\boxed{\rho^\star(\gamma, n) \longrightarrow \frac{1}{2}\log\left(1 + \frac{\gamma}{q}\right).} \tag{25}$$

In the large-system limit, local information captures the share $q/(q+\gamma)$ of decision variation; the remainder requires nonlocal state.

When $\gamma = 0$, the decisions decouple and $\eta_S(0) = 0$: current local information is sufficient, so any positive global delay makes the fresh-local architecture preferable. As γ increases, each agent's optimal action depends more strongly on remote states, $\eta_S(0)$ increases, and the crossover moves to larger τ/T. Thus stronger coordination coupling makes the system willing to tolerate increasingly stale-global information before abandoning global coordination in favor of fresh-local decisions.

4 Choosing the observation radius

We first ask when a wider synchronized snapshot is worth its extra delay, then compute the optimal radius in a concrete spatial model.

4.1 When is one more hop worthwhile?

Let D be the maximum allowed radius and let $\tau : [0, D] \to [0, \infty)$ be nondecreasing, with $\tau(0) = 0$. Define the fraction of open-loop loss eliminated by the snapshot as

$$G(r) = \underbrace{1 - \eta_S(r)}_{\text{share captured by a fresh snapshot}} \quad \underbrace{e^{-2\tau(r)/T}}_{\text{fraction surviving the delay}}. \tag{26}$$

Writing $R(r) = R(\mathcal{A}_{r,\tau(r)})$, Theorem 1 gives $R(r) = R_\infty[1 - G(r)]$, so the best radius maximizes $G(r)$.

Suppose a nearby sensor has already reported a new reading when an older reading arrives from farther away. To form a synchronized snapshot, the agent uses an older reading from the nearby sensor too. This is why widening the snapshot can hurt. An agent that keeps both readings can always ignore the distant one, so widening its access cannot increase its optimal decision cost.

On a finite graph, one more hop improves performance exactly when

$$G(r+1) > G(r). \tag{27}$$

Compare all admissible integer radii on a general graph; stopping at the first decline requires unimodality. Likewise, rounding a continuous optimum to a neighboring integer is justified only for a unimodal interpolated regret curve. Radius zero gives fresh-local information, while the diameter of a finite connected graph gives a complete stale-global snapshot.

The loss separates into the two components plotted in Fig. 2:

$$\frac{R(r)}{R_\infty} = \underbrace{1 - e^{-2\tau(r)/T}}_{\text{loss from delay}} + \underbrace{e^{-2\tau(r)/T}\eta_S(r)}_{\text{remaining loss from limited scope}}. \tag{28}$$

The delay term grows with radius, while the remaining scope term falls.

Proposition 2 (Continuous marginal-balance characterization)**.** *Under the assumptions of Theorem 1, additionally assume that $\tau(r)$ and $\eta_S(r)$ are continuously differentiable on $[0, D]$, with*

$$\tau'(r) \geq 0, \qquad \eta_S'(r) \leq 0, \qquad 0 \leq \eta_S(r) < 1.$$

Every interior optimal radius $r^\star \in (0, D)$ satisfies

$$\frac{d}{dr} \log[1 - \eta_S(r)]\bigg|_{r=r^\star} = \frac{2\tau'(r^\star)}{T}. \tag{29}$$

Proof. Since $G(r) > 0$, maximizing G is equivalent to maximizing $\log G$. At an interior optimum, $(\log G)' = 0$; differentiating Eq. (26) gives the stated equality. □

The proportional gain from wider scope equals the proportional loss from delay. If the difference between the left and right sides of Eq. (29), evaluated at radius r, is strictly decreasing in r, its sign identifies the unique optimum: radius zero if the difference is nonpositive at zero, radius D if it is nonnegative at D, and otherwise the unique interior crossing. Holding the spatial omission curve fixed and preserving this single-crossing condition, increasing T moves the preferred radius weakly outward, while increasing the latency scale in $\tau(r) = \alpha\tau_0(r)$ moves it weakly inward. The optimum can stay unchanged at a boundary. These conclusions concern decision-relevant prediction, rather than changes in raw state covariance.

4.2 An exponential omission curve

An exponential fresh omission curve gives a global radius solution using one derivative. This generic calculation will also apply to the spatial model in the next subsection.

Proposition 3 (Optimal radius for exponential omission)**.** *Under Theorem 1, suppose $\eta_S(r) = \eta_0 e^{-2r/\ell_c}$ with $0 < \eta_0 < 1$, $\ell_c > 0$, and $\tau(r) = r/v$ with $v > 0$. For continuous $r \geq 0$, the unique global optimal radius is*

$$\boxed{r^\star = \frac{\ell_c}{2}\left[\log\left(\eta_0\left(1 + \frac{vT}{\ell_c}\right)\right)\right]_+}, \tag{30}$$

where $[z]_+ = \max\{z, 0\}$. It is positive exactly when $\eta_0(1 + vT/\ell_c) > 1$. With a continuous radius cap D, $r_D^\star = \min\{D, r^\star\}$.

Proof. Differentiating the captured benefit gives

$$G'(r) = 2e^{-2r/(vT)}\left[\eta_0\left(\frac{1}{vT} + \frac{1}{\ell_c}\right)e^{-2r/\ell_c} - \frac{1}{vT}\right].$$

The bracket decreases strictly to a negative limit. If its value at zero is nonpositive, G decreases for all positive radii and zero is its unique maximum, including when the initial derivative is zero. Otherwise G increases up to its unique crossing and decreases thereafter. Solving for the crossing gives the formula; a cap truncates the maximizing radius. □

The cap in Fig. 2 limits observation of an infinite field, so information beyond D remains omitted. A graph diameter is different: the fresh view then includes the whole finite graph and omission is zero. The generic result assumes an omission curve; the next model derives it from state statistics and the desired action.

4.3 The canonical three-length result

We now specialize to a Gaussian spatial field on a dense one-dimensional network, represented by locations $z \in \mathbb{R}$. Let $\ell_s > 0$ be the spatial correlation length, the distance over which field values remain correlated, and let $T > 0$ be the temporal coherence time. The zero-mean field has covariance

$$\mathbb{E}[\theta(z,t)\theta(z',s)] = \sigma^2 e^{-|z-z'|/\ell_s} e^{-|t-s|/T}. \tag{31}$$

At a fixed time, this exponential covariance makes the field a spatial Ornstein–Uhlenbeck process [14]. Given the field value at a point, its two sides are conditionally independent; given an observed interval, each unobserved tail depends only on the adjacent boundary value. This spatial Markov property is used in Lemma 1. The decision-relevance length $\ell_c > 0$ determines how far remote states influence the desired local action. Define

$$x^\star(z,t) = \int_{-\infty}^{\infty} w_{\ell_c}(u)\theta(z+u,t)\,du, \qquad w_{\ell_c}(u) = \frac{1}{2\ell_c} e^{-|u|/\ell_c}. \tag{32}$$

The kernel integrates to one. Each location incurs squared tracking loss

$$f_z(x,\theta) = \frac{q}{2}[x(z) - x^\star(z,\theta)]^2, \qquad q > 0. \tag{33}$$

There is no action coupling, so the optimal policy is the conditional mean. We measure regret per unit length, equal by stationarity to expected pointwise loss, with $R_\infty = (q/2)\,\mathrm{Var}(x^\star(z,t))$. The radius-$r$ observation consists of all field values within distance r:

$$\mathcal{O}_r(z,t) = (\theta(z+u,t) : |u| \le r). \tag{34}$$

Information propagates at speed $v > 0$, giving

$$\tau(r) = r/v, \qquad L_T = vT. \tag{35}$$

The propagation length vT is the distance information travels in one coherence time.

Lemma 1 (Exponential omission by self-similarity)**.** *For the field* (31)*, decision* (32)*, and observation* (34)*, the normalized zero-delay spatial omission error is*

$$\eta_S(r) = \eta_0 e^{-2r/\ell_c}, \qquad \eta_0 = \frac{\ell_c}{\ell_c + 2\ell_s}. \tag{36}$$

Proof. By stationarity, set $z = 0$ and suppress time. The observed middle of the target integral contributes no uncertainty. Conditional on the two boundary values, the unobserved tails are independent by the spatial Markov property. Moving each tail outward by r scales its weight by e^{-r/ℓ_c}, while stationarity leaves its conditional variance relative to its boundary unchanged. Squaring the weight therefore gives $\eta_S(r) = e^{-2r/\ell_c}\eta_S(0)$. The covariance integrals in Appendix C give $\eta_S(0) = \ell_c/(\ell_c + 2\ell_s)$. □

Theorem 3 (Optimal synchronized-snapshot radius in the canonical model)**.** *Suppose the field obeys Eq.* (31)*, the target is Eq.* (32)*, the loss is Eq.* (33)*, and* $\tau(r) = r/v$*. The unique optimal radius on the infinite line is*

$$\boxed{r^\star = \frac{\ell_c}{2}\left[\log\frac{\ell_c + vT}{\ell_c + 2\ell_s}\right]_+}, \tag{37}$$

where $[z]_+ = \max\{z, 0\}$*. In particular,*

$$r^\star > 0 \quad \Longleftrightarrow \quad vT > 2\ell_s. \tag{38}$$

Within the interior regime, $r^\star$*:*

1. *increases with the temporal coherence time* T*;*
2. *increases with the information-propagation speed* v*;*
3. *decreases with the spatial correlation length* ℓ_s*; and*
4. *increases with the decision-relevance length* ℓ_c*.*

Proof. Lemma 1 gives $\eta_0 = \ell_c/(\ell_c + 2\ell_s)$. Substitution into Proposition 3 gives the radius and the threshold $vT > 2\ell_s$. In the positive-radius regime, $\partial r^\star/\partial(vT) = \ell_c/[2(\ell_c + vT)] > 0$ and $\partial r^\star/\partial \ell_s = -\ell_c/(\ell_c + 2\ell_s) < 0$. The integral representation in Appendix C proves the increase with ℓ_c. Here the omission prefactor varies with the lengths; these comparative statics do not hold automatically for the generic model with fixed η_0. □

Equation (30) also covers the free-η_0 curves in Fig. 2; its right panel reaches the widest allowed radius on an infinite field and still omits information beyond D. The threshold $vT > 2\ell_s$ has a direct interpretation: when nearby states are strongly correlated, the local reading already tells the agent more, so propagation must be faster before a wider view pays.

When $\ell_c \gg vT, \ell_s$ and $vT > 2\ell_s$, $r^\star \to (vT - 2\ell_s)/2$ as $\ell_c \to \infty$ with vT, ℓ_s fixed. Even as increasingly distant states matter to the decision, propagation delay limits how far the agent should look.

5 Numerical Illustrations

We ask when fresh local information beats a delayed global view, why an intermediate snapshot radius can be best, whether changing the decision alone can change the best radius, and how the framework applies across graph geometries. The examples use common-rate Gaussian-affine environments. Appendix D gives the settings and sampling procedures needed to reproduce the four figures.

5.1 When does fresh local information beat a delayed global view?

Stronger decision coupling increases the delay worth tolerating for a global view. Figure 1 illustrates this in the shared-resource model of Section 3.2. The exact finite-n boundary is Eq. (23); its large-n limit is $\rho^\star = \frac{1}{2}\log(1 + \gamma/q)$ from Eq. (25). Above the boundary, fresh local information gives lower regret. The finite-size curves approach the large-system limit as n increases.

We fix $q = 1$ and compare $n \in \{5, 10, 25, 100\}$ over nine coupling values. The Monte Carlo estimates infer the crossover from paired local and open-loop losses sampled at the current time, using $\widehat{\rho}^\star = \frac{1}{2}\log[\widehat{R}_\infty/(\widehat{R}_\infty - \widehat{R}_{\mathrm{loc}})]$. Their agreement with the exact curves checks the loss calculation and the inferred boundary; the estimates do not sample a separate delayed system at each age.

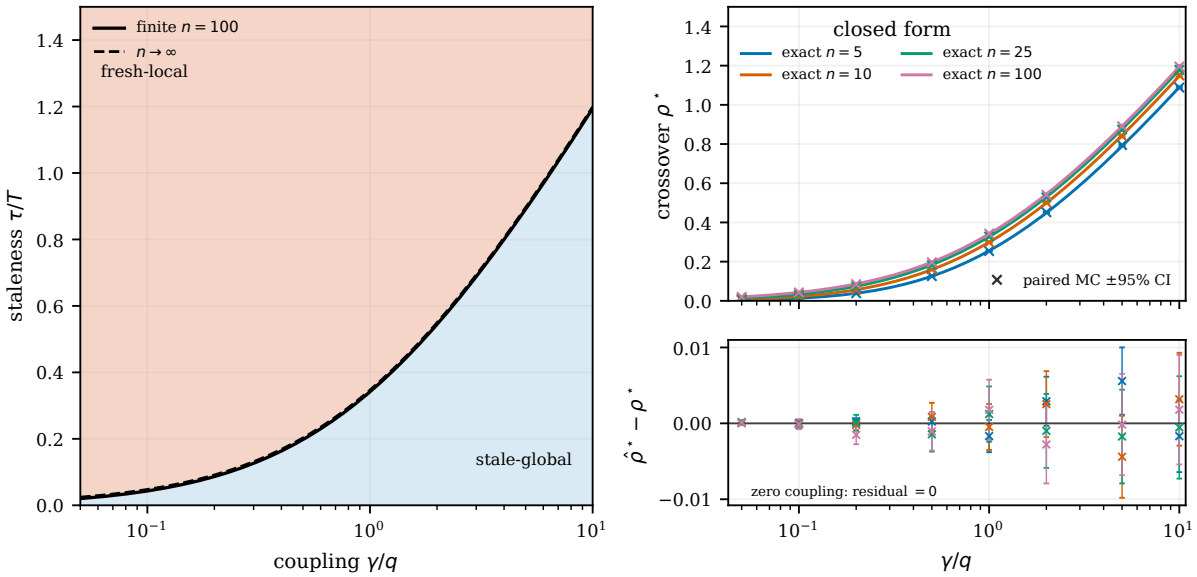


Figure 1: Fresh-local versus stale-global crossover in the shared-resource model. Left: the exact finite-$n = 100$ boundary (solid) and large-n limit (dashed). Upper right: exact boundaries and Monte Carlo estimates inferred from sampled local and open-loop losses for four system sizes, with paired delta-method 95% intervals. Lower right: estimated minus exact boundary. The zero-coupling boundary is zero and is omitted from logarithmic axes.

5.2 Why can an intermediate radius be best?

Figure 2 shows three possibilities for a synchronized snapshot: the best radius is local (left), interior (middle), or as wide as the imposed cap permits (right). With $\eta_S(r) = \eta_0 e^{-2r/\ell_c}$ and $\tau(r) = r/v$, the temporal loss grows with radius, while the residual spatial loss falls. Their sum can therefore be minimized at either endpoint or at an interior balance. For $\ell_c = 1$ and cap $D = 2$, the three settings give constrained analytical optima 0, 0.62638, and 2, respectively. The cap limits observation on an infinite spatial field: even at D, the snapshot omits some of that field.

The open markers check this decomposition using independent Gaussian samples with prescribed scalar component variances. They check variance addition in Eq. (28), rather than simulate a spatial field. The spatial and temporal curves explain which component limits the value of widening each snapshot.

5.3 Can the best radius change on the same network?

On the same 64-node ring, changing only the decision operator moves the optimal radius from 1 to 2 to 5 hops (Fig. 3). The graph, state covariance, coherence time $T = 6$, and latency 0.12 per hop stay fixed. With $Q = I$, we use the row-normalized operators $K_{ij} = e^{-d_G(i,j)/\ell_c} / \|e^{-d_G(i,\cdot)/\ell_c}\|_2$ for $\ell_c = 0.7, 2, 6$. The fixed covariance is

$$\Sigma_S = \sigma^2(\kappa_s^2 I + L_G)^{-\nu}, \qquad \kappa_s, \nu > 0, \tag{39}$$

where L_G is the ring Laplacian and the scale gives average marginal variance one. This covariance emphasizes low graph-frequency modes [15, 16].

Exact Gaussian conditioning gives each operator's omission curve, and the composition law gives regret at every integer radius. Each curve is normalized by its own $R_\infty = \frac{1}{2}\operatorname{tr}(K\Sigma_S K^\top)$. The

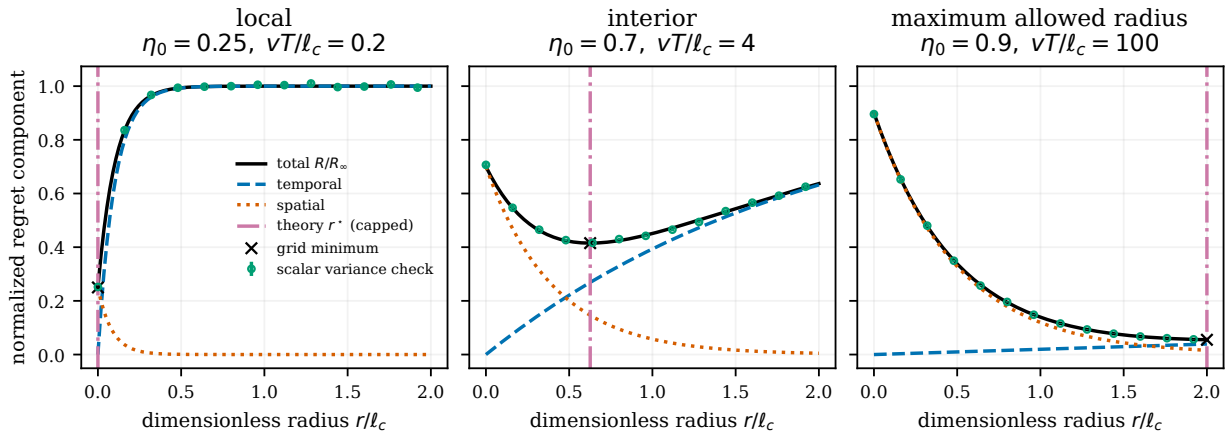


Figure 2: Local, interior, and capped optimal radii, from left to right. Curves show total regret and its temporal and residual-spatial components under exponential omission and linear delay. Open markers with 95% intervals use prescribed scalar variances. Dash-dotted lines and crosses mark analytical constrained and grid minimizers. At the cap, $\eta_S(D) > 0$: the snapshot still omits part of the infinite field.

raw state correlation is identical in all three cases, but the value of remote observations depends on how they affect the decision. State correlation alone therefore cannot determine the preferred radius. This ring example illustrates the exact one-hop comparison in Eq. (27).

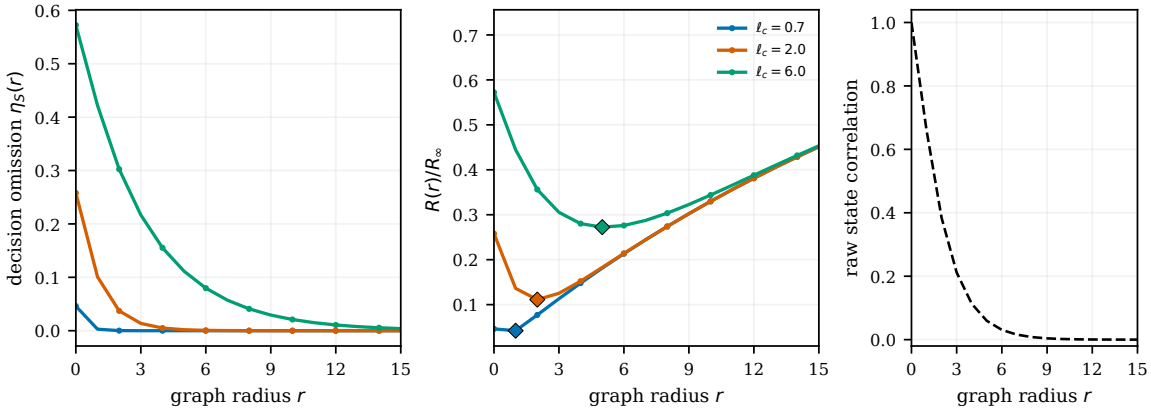


Figure 3: Changing the decision on a fixed 64-node ring. Left: decision-weighted omission for three decision operators. Middle: composed regret, normalized by each operator's open-loop regret, with optimal radii 1, 2, and 5 hops marked. Right: raw state correlation, identical across operators. The graph, covariance, temporal coherence, and per-hop latency are fixed.

5.4 Does the framework extend beyond the canonical spatial model?

The closed-form radius in Theorem 3 assumes a homogeneous infinite line. To illustrate the more general composition law of Theorem 1, we return to the finite-graph quadratic model and compare one 64-node path, ring, 8×8 grid, and random geometric graph. All use $Q = I$, the Laplacian covariance in Eq. (39) with $(\kappa_s, \nu) = (0.8, 1.5)$, and average marginal variance one. The reference

setting fixes $T = 5$, decision range $\ell_c = 2$, and delay $\tau(r) = 0.12r$.

For each graph, the row-normalized operator $K_{ij} \propto e^{-d_G(i,j)/\ell_c}$ defines the desired action. Gaussian conditional covariance gives each agent's residual action variance given its fresh neighborhood. Summing these variances and dividing by $\mathrm{tr}(K\Sigma_S K^\top)$ gives $\eta_S(r)$. The common-rate composition law then gives delayed regret. We evaluate every integer radius from zero through the graph diameter and select the minimum, without assuming unimodality or fitting an exponential omission curve.

Figure 4 gives reference optima of 2, 2, 3, and 2 hops, respectively: all are interior. Across the displayed decision-range sweep, the path and ring also have local optima at $\ell_c = 0.7$; none of the displayed settings selects the full graph. The sampled coherence-time and decision-range sweeps weakly increase the preferred radius on these four instances. These calculations illustrate architecture selection on finite graphs when no closed-form radius is available. They use synchronized snapshots and common-rate Gaussian dynamics; they neither extend the canonical radius formula to these geometries nor establish a new theorem.

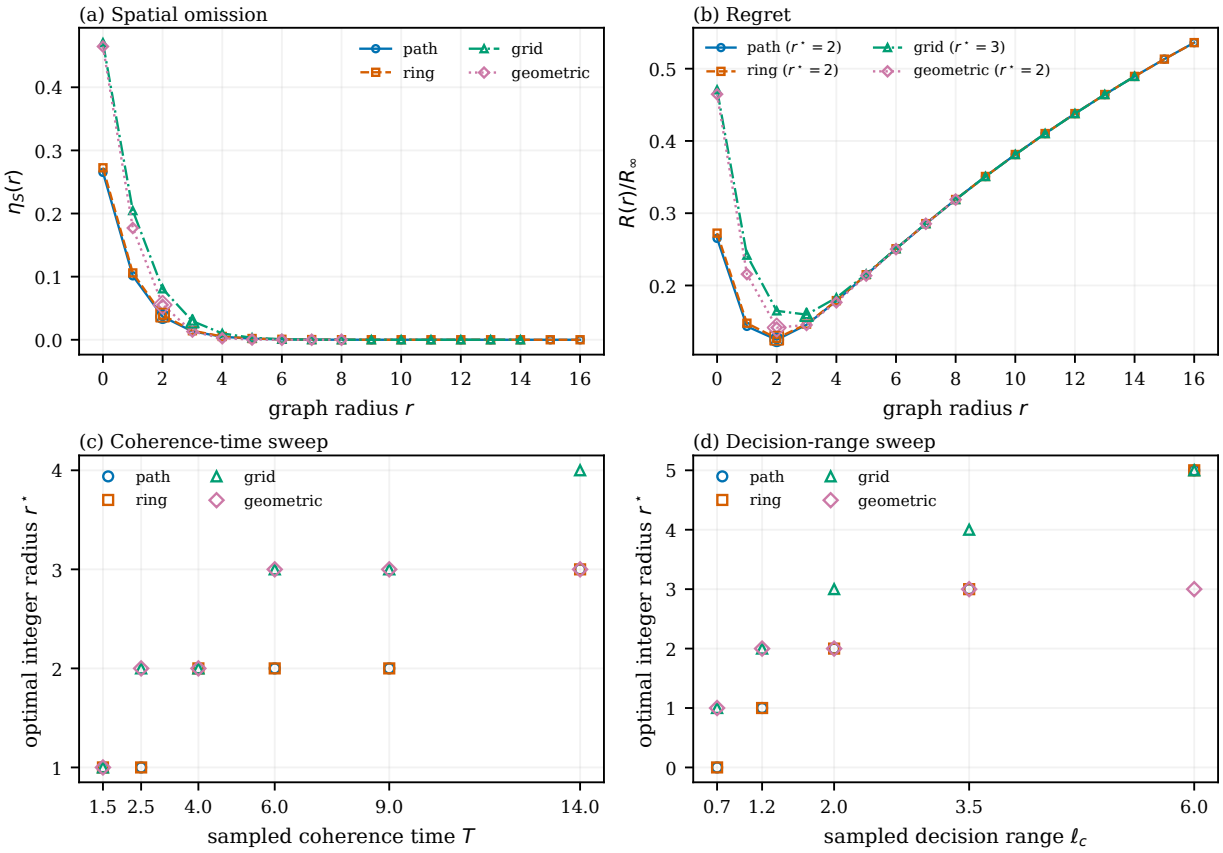


Figure 4: Composition-law evaluation on four 64-node graphs. Top: spatial omission and delayed regret with $Q = I$, $T = 5$, $\ell_c = 2$, and delay $0.12r$; curves show integer radii up to 16, or the diameter if smaller, while optimization uses every radius. Bottom: varying T or ℓ_c alone, holding each graph and covariance fixed. Markers show sampled integer optima, not transition locations. The canonical radius formula is not used.

6 Related Work

We organize the closest literatures around three distinctions: iterative computation versus information architecture, controlled dynamics versus an exogenous environment, and age-based freshness versus decision-relevant value. These distinctions frame the more specific connections developed below.

6.1 Static Team Decision Theory

Radner established foundational optimality conditions for static teams, in which decision makers share a common objective but act on different information [10]. Under appropriate convexity and regularity conditions, the resulting stationarity conditions are sufficient for global team optimality; in the quadratic Gaussian specialization, the optimal decision rules are affine [10, 11]. Lemma 2 specializes this framework to the weighted prediction problem induced by a deterministic quadratic Hessian; it is not a literal restatement of Radner's theorem. Krainak, Speyer, and Marcus subsequently relaxed the sufficient conditions associated with Radner's stationarity result and extended the theory to an exponential-quadratic criterion with jointly Gaussian state and observations [17]. In a companion paper, they formulated affine linear-quadratic-Gaussian and linear-exponential-Gaussian team problems as constrained parameter optimizations and represented optimal team gains as projections of the corresponding centralized gains [12].

Classical team theory also treated information itself as an organizational design object. Marschak and Radner analyzed the value of information, delayed information and the tradeoff between timeliness and completeness, and formulated team organization in terms of optimal networks [11]. More recently, Summers, Li, and Kamgarpour formulated information-structure design as the selection of additional measurements or communication links jointly with the corresponding team decision strategies [18]. Afshari and Mahajan characterized quadratic Gaussian static teams whose observations decompose into information common to all agents and agent-specific local information [19]. They subsequently applied this decomposition to team-optimal decentralized filtering over communication graphs with delayed information sharing, where sufficiently old measurements constitute common information while more recent measurements remain local [20]. Rusmevichientong and Van Roy studied large teams in which each agent observes the local cost structure only within a graph radius, using centralized performance as a benchmark [21]. Gamarnik, Goldberg, and Weber studied decentralized decisions based on local information in random decision networks and related their near-optimality to a correlation-decay, or long-range-independence, property [22].

Relative to these works, the distinctive object studied here is a structured family of information architectures in which spatial scope and temporal freshness are coupled through the scope–latency relation $\tau(r)$. Each spatial scope induces an acquisition delay, and the resulting architecture is valued through its prediction error for the full-information decision.

6.2 Distributed Optimization, Communication Constraints, and Scope

Distributed optimization studies how agents compute a common solution using local computation and message exchange. Classical asynchronous gradient methods allow delayed interprocessor communication while retaining convergence guarantees under suitable boundedness conditions [23], while distributed subgradient methods establish convergence when agents exchange information locally over time-varying network topologies [3]. In these canonical formulations, the communication process constrains an iterative computation of an optimizer. Our setting separates that computational question from information architecture: the full-information decision at each epoch is the benchmark, and the object of design is the information available when that decision must be made.

Networked and decentralized control place communication restrictions directly in the feedback loop, including sampling, packet loss, communication topology, and delay [24]. For spatially invariant plants, early work showed that optimal controllers can exhibit inherent spatial localization [5], while subsequent work imposed explicit finite-speed communication structure. Arbelaiz, Bamieh, Hosoi, and Jadbabaie characterize the spatial decay of optimal estimator gains, using the decay rate as a proxy for the distance over which measurements are useful in a spatially distributed dynamic estimation problem [25]. Voulgaris, Bianchini, and Bamieh studied H_2 control under delayed communication requirements [6]; Bamieh and Voulgaris formulated distance-dependent information propagation through funnel causality [7]; and Fardad and Jovanović developed state-space models for distributed controllers with finite communication speed [8]. These works are important antecedents to our use of spatial scope and propagation time, but they optimize dynamic feedback for controlled plants rather than information-constrained stage decisions in an exogenous environment.

Communication structure has also been treated as a design variable. Langbort and Gupta study interconnection topology when communication is assigned a topology-dependent cost [26], and Matni co-designs communication-delay structure and H_2 control using regularization [27]. In wireless consensus, Vanka, Gupta, and Haenggi explicitly trade communication radius against interference-induced delay and show that the preferred connectivity depends on network geometry [28]. Most directly, Ballotta, Jovanović, and Schenato parameterize control architectures by communication scope: each agent receives measurements from nodes within a chosen number of communication hops, with an architecture-dependent delay that increases with that scope. After optimizing the feedback gains for each scope, they show that sparse architectures can outperform all-to-all feedback [9]. Ballotta and Gupta obtain an analogous conclusion when maximizing consensus convergence speed under hop-dependent communication delay [29]. These results establish that scope-dependent delay can itself induce a sparse or intermediate optimal control architecture.

Our setting differs in the decision problem: an exogenous-state static team, with the current full-information action as the benchmark and an architecture valued by the prediction error it induces for that action. Ballotta, Arbelaiz, Gupta, Schenato, and Jovanović study a different but complementary question for spatially invariant dynamic plants: how delayed measurements from different spatial locations alter the spatial locality of the optimal feedback kernel [30]. Their results operate through closed-loop plant dynamics and controller localization, whereas ours operate through prediction of an exogenously varying full-information decision.

6.3 Freshness and Age of Information

Age of Information (AoI) formalizes the timeliness of status information through the age of the freshest received update, $\Delta(t) = t - u(t)$. Seminal work showed that minimizing AoI is distinct from maximizing throughput or minimizing packet delay and initiated a large literature on update generation, queueing, and scheduling [2,31]. Subsequent work generalized linear age to application-specific nonlinear age penalties [32]. Connections to remote estimation make the distinction between age and task loss explicit: for a Gauss–Markov Ornstein–Uhlenbeck process, estimation error under signal-agnostic sampling is a nonlinear function of age, whereas signal-aware sampling can exploit the realized estimation error and outperform age-optimal policies [13].

More recent work has pushed freshness metrics toward task-aware objectives. Shisher et al. showed that remote-inference loss can be nonmonotonic in age and designed communication policies around inference performance [33]; subsequent work considers correlated sources for which inference penalties depend jointly on multiple sources' ages [34]. Li et al. go further toward decision-aware freshness by jointly optimizing sampling and remote decision making under stale observations [35]. Those works optimize when or which updates reach a fixed remote receiver; we instead choose the

information architecture available to a decentralized team: what is observed, from what spatial scope, and at what age. Because broader scope itself incurs greater acquisition delay, spatial omission and temporal staleness are coupled by the architecture.

6.4 Centralization, Delegation, and Organizational Economics

A closely related organizational-economics literature asks how dispersed information, communication constraints, and decision rights determine organizational form. Aoki compares hierarchical and horizontal information structures when local conditions are uncertain and central monitoring or response is limited [36]. In the information-processing tradition, Radner studies organizations of limited-capacity processors trading off processing resources against decision delay [37]; Bolton and Dewatripont derive communication networks from processing and communication costs [38]; and Van Zandt studies real-time decentralized processing in which computational delay constrains the use of recent information [39]. Dessein and Santos likewise study adaptation to local information versus coordination under imperfect communication [40].

A complementary delegation literature makes decision rights endogenous when information and objectives are dispersed. Aghion and Tirole distinguish formal from real authority [41], while Dessein studies delegation as an alternative to strategic communication [42]. Most closely related to our centralization framing, Alonso, Dessein, and Matouschek compare centralized and decentralized coordination when managers privately observe local conditions [4]. Our model abstracts from incentive conflicts and endogenous information acquisition: agents share a common objective and the environment evolves exogenously. The tradeoff studied here is spatio-temporal: broader synchronized snapshots can improve coordination but arrive older.

6.5 Spatial Statistics and Graph Signal Processing

Spatial statistics provides the closest statistical antecedent to our spatial-predictability model. Kriging casts spatial prediction as an optimal linear prediction problem under spatial dependence [43]. For lattice data, Besag's conditional formulation represents spatial dependence through local Markov neighborhoods [44]; Lindgren, Rue, and Lindström later connect Matérn Gaussian fields to sparse Gaussian Markov random fields through stochastic partial differential equation representations [16]. Space–time statistics treats separability as an optional modeling restriction: Gneiting constructs broad classes of stationary nonseparable covariance functions [45]. These works provide the statistical context for our covariance-based spatial model and its separable and nonseparable distinction.

Graph signal processing extends analogous ideas to irregular domains. Shuman et al. use graph-Laplacian eigenvectors to define a graph Fourier domain and associate low graph frequencies with signals that vary slowly over the graph [15]. Graph-sampling theory asks when bandlimited graph signals can be exactly recovered from partial observations [46], while stationary graph signal processing gives spectral descriptions and Wiener-type estimators for random graph signals [47]. Here the prediction target is the full-information action, so a state pattern's importance depends on the decision operator as well as on its spatial smoothness.

7 Discussion and Conclusions

What the examples show. On the same 64-node ring in Fig. 3, changing only the decision target moves the preferred radius from 1 to 2 to 5 hops. The graph, state covariance, coherence time, and per-hop latency stay fixed. Graph distance determines what it takes to acquire information; the

decision determines whether that information is useful. This suggests selecting remote observations for their effect on the action, rather than expanding scope solely because remote states remain correlated. Figure 4 also illustrates the common-rate composition law on finite path, ring, grid, and geometric graphs beyond the canonical spatial model, with optimal radii found by exhaustive comparison.

Scope and limitations. The environment evolves exogenously: actions do not change its future state or the information available to other agents. When actions affect later observations, they may also signal information; the dynamic-team extension is discussed below. The principal family uses synchronized snapshots, while mixed-age and history-based architectures can retain information that these snapshots discard. Stationarity makes the expected snapshot regret time-invariant; Appendix A distinguishes this from fixed-epoch optimization and sample-path averages.

The projection identity assumes quadratic costs, a deterministic positive- definite Hessian, and unconstrained Euclidean actions. The explicit stale-global and composition laws additionally require affine full-information decisions and a stationary common-rate innovation structure with conditionally mean-zero innovations. Gaussianity is used to evaluate spatial prediction errors from covariances in the canonical model and the finite-graph examples in Sections 5.3 and 5.4; the canonical radius formula further uses a separable homogeneous field and linear propagation delay. Nonquadratic or discrete decisions, state-dependent Hessians, and feasibility constraints such as nonnegativity, capacity, or simplex restrictions require different analytical tools. Outside the stated assumptions, one should work directly with decision-relevant prediction-error functions or suitable bounds rather than assume the same factorization.

The numerical illustrations specialize the framework to Gaussian models. Heterogeneous temporal rates, nonseparable covariance, nonlinear desired actions, and mixed-age architectures remain directions for further study.

When spatial patterns change at different rates, a single coherence time can overstate or understate the value of delayed information, depending on which patterns drive the decision.

Most important extensions. A first extension is to price information acquisition explicitly. The present model represents communication and sensing constraints mainly through the feasible scope–delay relationship. Adding an acquisition cost or budget could produce sparse, nonuniform topologies and compressed summaries that preserve decision-relevant modes with less latency.

A second extension is adaptation to unknown or nonstationary statistics. Temporal coherence, spatial predictability, and decision sensitivity may need to be estimated online while the architecture changes on a slower timescale than operational decisions. The challenge is to account for estimation uncertainty while retaining a stable architecture-selection rule.

The most important extension is action-dependent state evolution and dynamic teams, where actions alter future state and may signal information to other agents. Extending information-architecture design beyond the exogenous-state baseline is therefore not a routine modification. Witsenhausen's counterexample [48] explains why: once control and signaling interact, the fixed-subspace projection characterization generally no longer solves the team problem.

Design implication. Taken together, the theory yields a practical design heuristic: expand information scope as long as the marginal decision value of broader information exceeds the predictive value lost in obtaining it.

AI Use Statement. The development of this paper made substantial use of AI tools, including Claude Code, ChatGPT/Codex, Gemini, and Grok, for drafting and revising text, developing code and figures, mathematical formulation and proofs. The human authors formulated the original problem, made substantial contributions throughout the paper's development, reviewed the manuscript, and accept full responsibility for its contents.

Code and Data Availability. Appendices B–D provide detailed algebraic derivations and numerical reproducibility information; Appendix A contains the projection proof and coupled optimality equations. A separate mathematical companion in the repository provides background on the projection and stochastic-process tools used here. The manuscript source, numerical code, experiment configurations, processed reference results, and figure-generation pipeline are available in the public TINA repository at `https://github.com/ANRGUSC/TINA` [49].

References

[1] M. Mitzenmacher, "How useful is old information?" *IEEE Transactions on Parallel and Distributed Systems*, vol. 11, no. 1, pp. 6–20, Jan. 2000.

[2] S. Kaul, R. Yates, and M. Gruteser, "Real-time status: How often should one update?" in *2012 Proceedings IEEE INFOCOM*. IEEE, 2012, pp. 2731–2735.

[3] A. Nedic and A. Ozdaglar, "Distributed subgradient methods for multi-agent optimization," *IEEE Transactions on Automatic Control*, vol. 54, no. 1, pp. 48–61, 2009.

[4] R. Alonso, W. Dessein, and N. Matouschek, "When does coordination require centralization?" *American Economic Review*, vol. 98, no. 1, pp. 145–179, 2008.

[5] B. Bamieh, F. Paganini, and M. A. Dahleh, "Distributed control of spatially invariant systems," *IEEE Transactions on Automatic Control*, vol. 47, no. 7, pp. 1091–1107, Jul. 2002.

[6] P. G. Voulgaris, G. Bianchini, and B. Bamieh, "Optimal H2 controllers for spatially invariant systems with delayed communication requirements," *Systems & Control Letters*, vol. 50, no. 5, pp. 347–361, Dec. 2003.

[7] B. Bamieh and P. G. Voulgaris, "A convex characterization of distributed control problems in spatially invariant systems with communication constraints," *Systems & Control Letters*, vol. 54, no. 6, pp. 575–583, Jun. 2005.

[8] M. Fardad and M. R. Jovanović, "Design of optimal controllers for spatially invariant systems with finite communication speed," *Automatica*, vol. 47, no. 5, pp. 880–889, May 2011.

[9] L. Ballotta, M. R. Jovanović, and L. Schenato, "Can decentralized control outperform centralized? the role of communication latency," *IEEE Transactions on Control of Network Systems*, vol. 10, no. 3, pp. 1629–1640, Sep. 2023.

[10] R. Radner, "Team decision problems," *The Annals of Mathematical Statistics*, vol. 33, no. 3, pp. 857–881, 1962.

[11] J. Marschak and R. Radner, *Economic Theory of Teams*, ser. Cowles Foundation Monographs. New Haven, CT: Yale University Press, 1972, no. 22.

[12] J. C. Krainak, J. Speyer, and S. I. Marcus, "Static team problems–part II: Affine control laws, projections, algorithms, and the LEGT problem," *IEEE Transactions on Automatic Control*, vol. 27, no. 4, pp. 848–859, 1982.

[13] T. Z. Ornee and Y. Sun, "Sampling and remote estimation for the Ornstein–Uhlenbeck process through queues: Age of information and beyond," *IEEE/ACM Transactions on Networking*, vol. 29, no. 5, pp. 1962–1975, 2021.

[14] C. E. Rasmussen and C. K. Williams, *Gaussian Processes for Machine Learning*. MIT press Cambridge, MA, 2006.

[15] D. I. Shuman, S. K. Narang, P. Frossard, A. Ortega, and P. Vandergheynst, "The emerging field of signal processing on graphs: Extending high-dimensional data analysis to networks and other irregular domains," *IEEE Signal Processing Magazine*, vol. 30, no. 3, pp. 83–98, 2013.

[16] F. Lindgren, H. Rue, and J. Lindström, "An explicit link between Gaussian fields and Gaussian Markov random fields: the stochastic partial differential equation approach," *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 73, no. 4, pp. 423–498, 2011.

[17] J. C. Krainak, J. Speyer, and S. I. Marcus, "Static team problems–part i: Sufficient conditions and the exponential cost criterion," *IEEE Transactions on Automatic Control*, vol. 27, no. 4, pp. 839–848, 1982.

[18] T. Summers, C. Li, and M. Kamgarpour, "Information structure design in team decision problems," *IFAC-PapersOnLine*, vol. 50, no. 1, pp. 2530–2535, 2017.

[19] M. Afshari and A. Mahajan, "Static teams with common information," *IFAC-PapersOnLine*, vol. 50, no. 1, pp. 11 926–11 931, 2017.

[20] ——, "Team optimal decentralized state estimation," in *2018 IEEE Conference on Decision and Control (CDC)*. IEEE, 2018, pp. 5044–5050.

[21] P. Rusmevichientong and B. Van Roy, "Decentralized decision-making in a large team with local information," *Games and Economic Behavior*, vol. 43, no. 2, pp. 266–295, 2003.

[22] D. Gamarnik, D. A. Goldberg, and T. Weber, "Correlation decay in random decision networks," *Mathematics of Operations Research*, vol. 39, no. 2, pp. 229–261, 2014.

[23] J. N. Tsitsiklis, D. P. Bertsekas, and M. Athans, "Distributed asynchronous deterministic and stochastic gradient optimization algorithms," *IEEE Transactions on Automatic Control*, vol. 31, no. 9, pp. 803–812, Sep. 1986.

[24] J. P. Hespanha, P. Naghshtabrizi, and Y. Xu, "A survey of recent results in networked control systems," *Proceedings of the IEEE*, vol. 95, no. 1, pp. 138–162, Jan. 2007.

[25] J. Arbelaiz, B. Bamieh, A. E. Hosoi, and A. Jadbabaie, "Optimal estimation in spatially distributed systems: How far to share measurements from?" *IEEE Transactions on Automatic Control*, vol. 70, no. 5, pp. 3226–3239, 2025. [Online]. Available: https://doi.org/10.1109/TAC.2024.3504257

[26] C. Langbort and V. Gupta, "Minimal interconnection topology in distributed control design," *SIAM Journal on Control and Optimization*, vol. 48, no. 1, pp. 397–413, 2009.

[27] N. Matni, "Communication delay co-design in $\mathcal{H}_2$-distributed control using atomic norm minimization," *IEEE Transactions on Control of Network Systems*, vol. 4, no. 2, pp. 267–278, Jun. 2017.

[28] S. Vanka, V. Gupta, and M. Haenggi, "Power-delay analysis of consensus algorithms on wireless networks with interference," *International Journal of Systems, Control and Communications*, vol. 2, no. 1/2/3, pp. 256–274, 2010.

[29] L. Ballotta and V. Gupta, "Faster consensus via a sparser controller," *IEEE Control Systems Letters*, vol. 7, pp. 1459–1464, 2023.

[30] L. Ballotta, J. Arbelaiz, V. Gupta, L. Schenato, and M. R. Jovanović, "The role of communication delays in the optimal control of spatially invariant systems," *IEEE Transactions on Automatic Control*, vol. 71, no. 2, pp. 978–993, Feb. 2026.

[31] R. D. Yates, Y. Sun, D. R. Brown, S. K. Kaul, E. Modiano, and S. Ulukus, "Age of information: An introduction and survey," *IEEE Journal on Selected Areas in Communications*, vol. 39, no. 5, pp. 1183–1210, 2021.

[32] Y. Sun, E. Uysal-Biyikoglu, R. D. Yates, C. E. Koksal, and N. B. Shroff, "Update or wait: How to keep your data fresh," *IEEE Transactions on Information Theory*, vol. 63, no. 11, pp. 7492–7508, 2017.

[33] M. K. C. Shisher, Y. Sun, and I.-H. Hou, "Timely communications for remote inference," *IEEE/ACM Transactions on Networking*, vol. 32, no. 5, pp. 3824–3839, 2024.

[34] M. K. C. Shisher, V. Tripathi, M. Chiang, and C. G. Brinton, "AoI-based scheduling of correlated sources for timely inference," *IEEE Transactions on Networking*, vol. 34, pp. 2181–2195, 2026.

[35] A. Li, S. Wu, G. C. F. Lee, and S. Sun, "From freshness to effectiveness: Goal-oriented sampling for remote decision making," *IEEE Transactions on Information Theory*, vol. 72, no. 4, pp. 2237–2276, Apr. 2026. [Online]. Available: https://doi.org/10.1109/TIT.2026.3663678

[36] M. Aoki, "Horizontal vs. vertical information structure of the firm," *The American Economic Review*, vol. 76, no. 5, pp. 971–983, 1986.

[37] R. Radner, "The organization of decentralized information processing," *Econometrica: Journal of the Econometric Society*, vol. 61, no. 5, pp. 1109–1146, 1993.

[38] P. Bolton and M. Dewatripont, "The firm as a communication network," *The Quarterly Journal of Economics*, vol. 109, no. 4, pp. 809–839, 1994.

[39] T. Van Zandt, "Real-time decentralized information processing as a model of organizations with boundedly rational agents," *The Review of Economic Studies*, vol. 66, no. 3, pp. 633–658, 1999.

[40] W. Dessein and T. Santos, "Adaptive organizations," *Journal of Political Economy*, vol. 114, no. 5, pp. 956–995, 2006.

[41] P. Aghion and J. Tirole, "Formal and real authority in organizations," *Journal of Political Economy*, vol. 105, no. 1, pp. 1–29, 1997.

[42] W. Dessein, "Authority and communication in organizations," *The Review of Economic Studies*, vol. 69, no. 4, pp. 811–838, 2002.

[43] N. Cressie, "The origins of kriging," *Mathematical Geology*, vol. 22, no. 3, pp. 239–252, 1990.

[44] J. Besag, "Spatial interaction and the statistical analysis of lattice systems," *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 36, no. 2, pp. 192–225, 1974.

[45] T. Gneiting, "Nonseparable, stationary covariance functions for space–time data," *Journal of the American Statistical Association*, vol. 97, no. 458, pp. 590–600, 2002.

[46] S. Chen, R. Varma, A. Sandryhaila, and J. Kovačević, "Discrete signal processing on graphs: Sampling theory," *IEEE Transactions on Signal Processing*, vol. 63, no. 24, pp. 6510–6523, 2015.

[47] N. Perraudin and P. Vandergheynst, "Stationary signal processing on graphs," *IEEE Transactions on Signal Processing*, vol. 65, no. 13, pp. 3462–3477, Jul. 2017.

[48] H. S. Witsenhausen, "A counterexample in stochastic optimum control," *SIAM Journal on Control*, vol. 6, no. 1, pp. 131–147, 1968.

[49] S. Moeller and B. Krishnamachari, "TINA: Reproduction package for "a theory of information architecture for networked decisions"," GitHub repository, 2026, source code, data, figures, and manuscript materials. [Online]. Available: https://github.com/ANRGUSC/TINA

A General formulation and fixed-architecture optimum

This appendix preserves the broader fixed-epoch quadratic formulation. Its target may be non-affine, unlike the temporal composition theorem.

On a probability space $(\Omega, \mathcal{F}, \mathbb{P})$, let $\theta = \{\theta_t\}_{t\in\mathbb{T}}$ be an environment process, with discrete or continuous time. Write

$$\theta_t = (\theta_{1,t}, \ldots, \theta_{n,t}) \in \mathbb{R}^m, \quad \theta_{i,t} \in \mathbb{R}^{m_i}, \quad \sum_i m_i = m, \quad \mathbb{E}\|\theta_t\|^2 < \infty. \tag{A.1}$$

Write $\bar{\theta} = \mathbb{E}[\theta_t]$ and $\Sigma_S = \mathrm{Cov}(\theta_t)$ for its mean and covariance at the fixed epoch. The local state $\theta_{i,t}$ may describe channel conditions, demand, workload, or sensing quality. This fixed-epoch formulation requires no stationarity, Gaussianity, independence, or Markov assumptions. Agent i chooses $x_{i,t} \in \mathbb{R}^{d_i}$; the joint action is $x_t = (x_{1,t}, \ldots, x_{n,t}) \in \mathbb{R}^d$, with $d = \sum_i d_i$.

The agents share a common per-stage cost

$$f(x, \theta) = \frac{1}{2} x^\top Q x - b(\theta)^\top x + c(\theta), \tag{A.2}$$

where $Q \in \mathbb{R}^{d\times d}$ is symmetric and positive definite, $b : \mathbb{R}^m \to \mathbb{R}^d$ is measurable, and $c : \mathbb{R}^m \to \mathbb{R}$ is an additive state-dependent term. We assume $\mathbb{E}\|b(\theta_t)\|^2 < \infty$ and $\mathbb{E}|c(\theta_t)| < \infty$. The off-diagonal blocks of Q couple agents' actions, while $b(\theta)$ determines which states affect the preferred action. This decision coupling need not coincide with proximity in G.

If the complete current state θ_t is available without information constraints, the unique optimal joint action is

$$x_t^\star \triangleq x^\star(\theta_t) = Q^{-1} b(\theta_t). \tag{A.3}$$

We call $x_t^\star$ the *clairvoyant* or *full-information* decision.

Completing the square in (A.2) gives

$$f(x, \theta) - f(x^\star(\theta), \theta) = \frac{1}{2} \left(x - x^\star(\theta)\right)^\top Q \left(x - x^\star(\theta)\right). \tag{A.4}$$

Hence, for the quadratic model, excess decision cost is exactly a weighted squared error in reproducing the full-information action.

For the affine specializations used to obtain closed forms, we write

$$b(\theta) = B\theta + b_0, \qquad B \in \mathbb{R}^{d\times m}, \qquad b_0 \in \mathbb{R}^d, \qquad K \triangleq Q^{-1}B \in \mathbb{R}^{d\times m}. \tag{A.5}$$

The matrix B maps state into the linear term of the objective, while K is the resulting decision-sensitivity operator. The general projection results do not require this affine form.

Centering the affine model. If $b(\theta) = B\theta + b_0$, write $\theta^c = \theta - \mathbb{E}[\theta]$ and $x^c = x - \mathbb{E}[x^\star]$. Then the centered desired action is $x^{\star,c} = K\theta^c$, with K as in Eq. (A.5). Translating each observation and each action by its known mean gives a bijection of the admissible rules and leaves $x - x^\star$ unchanged. Thus Eq. (A.4) is precisely the tracking loss of the main text in centered coordinates. For a non-affine b, the projection result below still holds, but the scalar state prediction assumption need not give scalar prediction of $x^\star$.

Information and admissible rules. Let $\mathcal{G}_i \subseteq \mathcal{F}$ be the sigma-algebra generated by agent i's observations; for example, $\mathcal{G}_i = \sigma(Z_i)$. Measurability here means that u_i can be computed from those observations. An architecture and its policy space are

$$\mathcal{A} = (\mathcal{G}_1, \dots, \mathcal{G}_n), \qquad \mathcal{S}(\mathcal{A}) = \{u : \mathbb{E}\|u\|^2 < \infty,\ u_i \text{ is } \mathcal{G}_i\text{-measurable for every } i\}. \tag{A.6}$$

In the local example $\mathcal{G}_{i,t}^{\mathrm{loc}} = \sigma(\theta_{i,t})$. If observations are enlarged without discarding old ones, then

$$\mathcal{G}_i \subseteq \mathcal{G}_i' \text{ (all } i) \quad \Longrightarrow \quad \mathcal{S}(\mathcal{A}) \subseteq \mathcal{S}(\mathcal{A}'), \quad R(\mathcal{A}') \le R(\mathcal{A}). \tag{A.7}$$

This formalizes the information-inclusion argument in the body.

Let

$$\mathcal{H} = L^2(\Omega, \mathcal{F}, \mathbb{P}; \mathbb{R}^d) \tag{A.8}$$

and equip $\mathcal{H}$ with the weighted inner product

$$\langle u, v\rangle_Q \triangleq \mathbb{E}\left[u^\top Q v\right], \qquad \|u\|_Q^2 \triangleq \langle u, u\rangle_Q. \tag{A.9}$$

Since $Q \succ 0$, the induced norm is equivalent to the ordinary L^2 norm.

Lemma 2 (Fixed-architecture quadratic projection; Radner-type specialization)**.** *For any static information architecture $\mathcal{A}$, the admissible policy set $\mathcal{S}(\mathcal{A})$ defined in* (A.6) *is a closed linear subspace of $\mathcal{H}$. Under the quadratic cost* (A.2)*, the unique optimal policy, up to almost-sure equality, implementable under $\mathcal{A}$ is*

$$u_{\mathcal{A}}^\star = \Pi_{\mathcal{S}(\mathcal{A})} x_t^\star, \tag{A.10}$$

where $\Pi_{\mathcal{S}(\mathcal{A})}$ denotes orthogonal projection with respect to (A.9)*. Consequently,*

$$\boxed{R(\mathcal{A}) = \frac{1}{2}\left\|x_t^\star - \Pi_{\mathcal{S}(\mathcal{A})} x_t^\star\right\|_Q^2.} \tag{A.11}$$

This identity specializes the classical static-team formulation [10] to the quadratic model used here.

Proof of Lemma 2.

Proof. Linearity of $\mathcal{S}(\mathcal{A})$ follows because $\mathcal{G}_i$-measurability is preserved under linear combinations. Closedness follows because an L^2 limit of $\mathcal{G}_i$-measurable random variables admits a $\mathcal{G}_i$-measurable version.

For any admissible policy u, (A.4) gives

$$\begin{aligned}\mathbb{E}\left[f(u, \theta_t) - f(x_t^\star, \theta_t)\right] &= \frac{1}{2}\mathbb{E}\left[(u - x_t^\star)^\top Q (u - x_t^\star)\right] \\ &= \frac{1}{2}\|u - x_t^\star\|_Q^2.\end{aligned} \tag{A.12}$$

Minimizing this expression over the closed subspace $\mathcal{S}(\mathcal{A})$ is the standard Hilbert-space projection problem, yielding (A.10) and (A.11). □

The projection error is Q-orthogonal to every admissible policy. In particular, $2R(\mathcal{A}) = \|x_t^\star\|_Q^2 - \|u_{\mathcal{A}}^\star\|_Q^2$, which is often useful for computing regret.

Corollary 2 (Coupled conditional normal equations)**.** *The unique team optimum in Lemma 2 satisfies*

$$\boxed{\mathbb{E}\big[\big(Qu^\star_{\mathcal{A}} - b(\theta_t)\big)_i \mid \mathcal{G}_i\big] = 0, \qquad i = 1, \dots, n.} \tag{A.13}$$

With Q partitioned into action blocks, this is equivalently

$$Q_{ii}u_i^\star + \sum_{j \neq i} Q_{ij}\,\mathbb{E}[u_j^\star \mid \mathcal{G}_i] = \mathbb{E}[b_i(\theta_t) \mid \mathcal{G}_i]. \tag{A.14}$$

These are the coupled person-by-person stationarity equations associated with the classical static-team formulation [10]. Strict convexity makes them sufficient for the unique team optimum here. With heterogeneous information and off-diagonal Q, one generally does not have $u_i^\star = \mathbb{E}[x_i^\star \mid \mathcal{G}_i]$. When all components share a common information sigma-algebra, deterministic Q makes the joint solution the ordinary vector conditional expectation.

Justification of the conditional normal equations. For an admissible perturbation h, differentiating expected quadratic cost at the optimum gives $\mathbb{E}[h^\top(Qu^\star - b(\theta_t))] = 0$. Choose perturbations supported on one agent's action block, with that block an arbitrary square-integrable function of its observations. Conditional expectation then gives Eq. (A.13). Conversely, these conditional equations make the linear term vanish in

$$\mathbb{E}[f(u^\star + h, \theta_t) - f(u^\star, \theta_t)] = \mathbb{E}[h^\top(Qu^\star - b(\theta_t))] + \tfrac{1}{2}\mathbb{E}[h^\top Q h].$$

Positive definiteness makes the remaining term positive for any nonzero L^2 perturbation, proving sufficiency and uniqueness up to almost-sure equality. □

Common information and degenerate normalization. If every agent observes the same Z, let $m = \mathbb{E}[x^\star \mid Z]$. For every admissible perturbation h, conditioning gives $\mathbb{E}[h^\top Q(x^\star - m)] = 0$. The error expansion then proves that m is the unique optimum, regardless of the off-diagonal entries of Q. For different observations this argument does not apply componentwise. More generally, a common old snapshot gives regret $\frac{1}{2}\mathbb{E}[(x_t^\star - \mathbb{E}[x_t^\star \mid \theta_{t-\tau}])^\top Q(x_t^\star - \mathbb{E}[x_t^\star \mid \theta_{t-\tau}])]$ even for a non-affine target. If its conditional mean converges in mean square to the unconditional mean, this regret converges to the no-information value. Under a stationary Markov environment, data processing makes this prediction loss nondecreasing in delay; arbitrary snapshot information sets themselves need not be nested. The best constant is $\mathbb{E}[x^\star]$; its regret is half the weighted variance of $x^\star$. In centered coordinates, $R_\infty = 0$ and $Q \succ 0$ imply $x_t^\star = 0$ almost surely. The zero rule has zero regret under every architecture, so the unnormalized composition law remains valid while normalized shares are undefined.

Stationarity makes expected snapshot regret time-invariant. It is not needed for the fixed-epoch projection result. Identifying long-run sample-path averages with these expectations additionally requires ergodicity.

A.1 Temporal special cases

The stationary Gaussian Ornstein–Uhlenbeck process is a sufficient example of the main conditional-mean assumption, rather than its definition:

$$d\theta_t = -\frac{1}{T}(\theta_t - \bar{\theta})\,dt + \sqrt{2/T}\,\Sigma_S^{1/2} dW_t. \tag{A.15}$$

Here W_t is standard multivariate Brownian motion and $\Sigma_S \succeq 0$. Its centered transition has independent Gaussian innovations, and

$$\mathrm{Cov}(\theta_t, \theta_s) = e^{-|t-s|/T}\Sigma_S. \tag{A.16}$$

For a stationary AR(1) process with coefficient $0 < c < 1$ and independent, mean-zero square-integrable innovations, the lag-k prediction coefficient is c^k and the composition weights are $1 - c^{2k}$ and c^{2k}. Neither statement asserts that a latest neighborhood snapshot is sufficient relative to its entire observation history.

B Aggregate-coupling calculations

This section supplies the matrix arithmetic for Proposition 1 above. The individual regrets are

$$R_\infty = \frac{n\sigma^2\,(q + \gamma(n-1))}{2q(q+\gamma n)}, \tag{B.1}$$

$$R_{\mathrm{loc}} = \frac{n\sigma^2\gamma^2(n-1)}{2q(q+\gamma)(q+\gamma n)}. \tag{B.2}$$

For $Q = qI + \gamma \mathbf{1}\mathbf{1}^\top$,

$$Q^{-1} = \frac{1}{q}\left(I - \frac{\gamma}{q+\gamma n}\mathbf{1}\mathbf{1}^\top\right).$$

Proof. Recall the coupled normal equations (A.14) from Appendix A:

$$Q_{ii}u_i + \sum_{j\neq i} Q_{ij}\mathbb{E}[u_j \mid \mathcal{G}_i] = \mathbb{E}[b_i(\theta_t) \mid \mathcal{G}_i].$$

Here $Q_{ii} = q + \gamma$, $Q_{ij} = \gamma$ for $j \neq i$, and $b_i(\theta) = \theta_i$. Let $\mu_j = \mathbb{E}[u_j]$. For $j \neq i$, the policy u_j is a function of $\theta_{j,t}$, which is independent of $\theta_{i,t}$, so

$$\mathbb{E}[u_j \mid \mathcal{G}^{\mathrm{loc}}_{i,t}] = \mathbb{E}[u_j] = \mu_j,$$

while $\theta_{i,t}$ is $\mathcal{G}^{\mathrm{loc}}_{i,t}$-measurable. The conditional normal equation (A.14) therefore gives

$$(q+\gamma)u_i + \gamma\sum_{j\neq i}\mu_j - \theta_{i,t} = 0.$$

Taking expectations yields

$$q\mu_i + \gamma\sum_j \mu_j = 0.$$

Summing over i implies $\sum_j \mu_j = 0$, hence every $\mu_i = 0$ and $u_i = \theta_{i,t}/(q+\gamma)$, which proves (20).

Open-loop regret. Here $B = I$ and $\Sigma_S = \sigma^2 I$, so by (6), $R_\infty = \frac{\sigma^2}{2}\,\mathrm{tr}(Q^{-1})$. From the Sherman–Morrison formula above,

$$\mathrm{tr}(Q^{-1}) = \frac{1}{q}\Big(n - \frac{\gamma n}{q+\gamma n}\Big) = \frac{n\,(q+\gamma(n-1))}{q\,(q+\gamma n)},$$

which gives (B.1).

Local regret. By the Pythagorean form of Lemma 2 (the orthogonality identity in Appendix A, with both $x^\star$ and u^{loc} mean zero), $2R_{\mathrm{loc}} = \|x^\star\|_Q^2 - \|u^{\mathrm{loc}}\|_Q^2 = 2R_\infty - \|u^{\mathrm{loc}}\|_Q^2$. Since $u^{\mathrm{loc}} = \theta/(q+\gamma)$ and $\mathrm{tr}(Q) = n(q+\gamma)$,

$$\|u^{\mathrm{loc}}\|_Q^2 = \frac{\sigma^2 \,\mathrm{tr}(Q)}{(q+\gamma)^2} = \frac{n\sigma^2}{q+\gamma}.$$

Therefore

$$R_{\mathrm{loc}} = \frac{n\sigma^2}{2}\left[\frac{q+\gamma(n-1)}{q(q+\gamma n)} - \frac{1}{q+\gamma}\right] = \frac{n\sigma^2}{2}\cdot\frac{(q+\gamma(n-1))(q+\gamma) - q(q+\gamma n)}{q(q+\gamma)(q+\gamma n)}.$$

The numerator reduces to $\gamma^2(n-1)$, which gives (B.2).

Share. Dividing (B.2) by (B.1), the factors $n\sigma^2/2$ and $q(q+\gamma n)$ cancel, leaving (21).

□

C Canonical spatial omission calculation

We expand the proof of Lemma 1 above, including the equal-inverse-length case in the variance integral. We use a continuous version of the field; observations on an interval can equivalently be generated by a countable dense subset (including the center at radius zero).

Proof. By spatial stationarity, take $z = 0$ and suppress the time argument. Write $\beta_c \triangleq 1/\ell_c$ and $\beta_s \triangleq 1/\ell_s$ for the inverse lengths, so $w(u) = \frac{\beta_c}{2}e^{-\beta_c|u|}$ and $\mathrm{Cov}(\theta(u),\theta(u')) = \sigma^2 e^{-\beta_s|u-u'|}$.

Step 0 (omission is a variance ratio). With $Q = qI$ the optimal policy given $\mathcal{O}_r$ is $\mathbb{E}[x^\star(0) \mid \mathcal{O}_r]$. Its expected loss is $\frac{q}{2}\,\mathbb{E}\big[\mathrm{Var}(x^\star(0) \mid \mathcal{O}_r)\big]$. For a Gaussian field the conditional variance is a deterministic number: it does not depend on the observed values (the Gaussian conditional covariance is independent of the observed values). Since $R_\infty = \frac{q}{2}\mathrm{Var}(x^\star(0))$,

$$\eta_S(r) = \frac{\mathrm{Var}(x^\star(0) \mid \mathcal{O}_r)}{\mathrm{Var}(x^\star(0))}.$$

Step 1 (split into observed middle and two tails). Write $x^\star(0) = M_r + T_r^+ + T_r^-$, where

$$M_r = \int_{-r}^{r} w(u)\theta(u)\,du, \qquad T_r^+ = \int_r^\infty w(u)\theta(u)\,du, \qquad T_r^- = \int_{-\infty}^{-r} w(u)\theta(u)\,du.$$

The middle M_r is computed from observed values, so it contributes no conditional variance.

Step 2 (rewrite the right tail). Substituting $u = r + s$,

$$T_r^+ = \frac{\beta_c}{2}e^{-\beta_c r}\,Y_r, \qquad Y_r \triangleq \int_0^\infty e^{-\beta_c s}\,\theta(r+s)\,ds.$$

Step 3 (spatial Markov property). By the exponential-covariance model in Section 4.3, the field is Markov in z. Given the observed segment $[-r, r]$, the right tail $\{\theta(u) : u > r\}$ and the left tail $\{\theta(u) : u < -r\}$ are conditionally independent, and each depends on the segment only through its adjacent endpoint $\theta(r)$ or $\theta(-r)$. Therefore

$$\mathrm{Var}(T_r^+ \mid \mathcal{O}_r) = \mathrm{Var}(T_r^+ \mid \theta(r)), \qquad \mathrm{Var}(T_r^- \mid \mathcal{O}_r) = \mathrm{Var}(T_r^- \mid \theta(-r)), \qquad \mathrm{Cov}(T_r^+, T_r^- \mid \mathcal{O}_r) = 0.$$

Step 4 (the tail variance does not depend on r). Define $C \triangleq \mathrm{Var}(Y_0 \mid \theta(0))$. Shifting the field by r does not change its law, so $\mathrm{Var}(Y_r \mid \theta(r)) = C$ for every r. By reflection symmetry (the covariance depends only on $|u - u'|$), the left tail gives the same value. Hence

$$\mathrm{Var}(x^\star(0) \mid \mathcal{O}_r) = 2\Big(\frac{\beta_c}{2}\Big)^2 e^{-2\beta_c r}\, C = \frac{\beta_c^2}{2} e^{-2\beta_c r}\, C.$$

At $r = 0$ the same argument applies with $\mathcal{O}_0 = \sigma(\theta(0))$ and $M_0 = 0$, giving $\mathrm{Var}(x^\star(0) \mid \theta(0)) = \frac{\beta_c^2}{2} C$. Dividing, $\eta_S(r) = \eta_S(0)\, e^{-2\beta_c r} = \eta_S(0)\, e^{-2r/\ell_c}$. The unknown constant C cancels, so it never has to be computed.

Step 5 (the constant $\eta_S(0)$). Conditioning on the single Gaussian variable $\theta(0)$ gives $\mathrm{Var}(x^\star(0) \mid \theta(0)) = \mathrm{Var}(x^\star(0))\big(1 - \mathrm{Corr}(x^\star(0), \theta(0))^2\big)$ (the Gaussian conditional covariance is independent of the observed values), so $\eta_S(0) = 1 - \mathrm{Corr}^2$. Two moments are needed.

Covariance with the local value:

$$\mathrm{Cov}(x^\star(0), \theta(0)) = \int_{\mathbb{R}} w(u)\, \sigma^2 e^{-\beta_s |u|}\, du = \frac{\sigma^2 \beta_c}{2} \int_{\mathbb{R}} e^{-(\beta_c + \beta_s)|u|}\, du = \frac{\sigma^2 \beta_c}{\beta_c + \beta_s}.$$

Variance of the target: $\mathrm{Var}(x^\star(0)) = \frac{\sigma^2 \beta_c^2}{4}\, I$, where $I \triangleq \iint e^{-\beta_c(|u| + |v|) - \beta_s |u - v|}\, du\, dv$.

To evaluate I, define the inner integral $h(u) \triangleq \int_{\mathbb{R}} e^{-\beta_c |v|} e^{-\beta_s |u - v|}\, dv$. Then h is even, so $I = 2\int_0^\infty e^{-\beta_c u} h(u)\, du$. For $u \geq 0$, split the v-integral at 0 and u:

$$h(u) = \underbrace{\frac{e^{-\beta_s u}}{\beta_c + \beta_s}}_{v<0} + \underbrace{\frac{e^{-\beta_c u} - e^{-\beta_s u}}{\beta_s - \beta_c}}_{0 \leq v \leq u} + \underbrace{\frac{e^{-\beta_c u}}{\beta_c + \beta_s}}_{v>u}.$$

(The middle term is read as $u\, e^{-\beta_c u}$ when $\beta_s = \beta_c$.) Multiply by $e^{-\beta_c u}$ and integrate over $u \geq 0$ term by term:

$$\frac{1}{(\beta_c + \beta_s)^2}, \qquad \frac{1}{\beta_s - \beta_c}\Big(\frac{1}{2\beta_c} - \frac{1}{\beta_c + \beta_s}\Big) = \frac{1}{2\beta_c(\beta_c + \beta_s)}, \qquad \frac{1}{2\beta_c(\beta_c + \beta_s)}.$$

Summing and doubling,

$$I = 2\left[\frac{1}{(\beta_c + \beta_s)^2} + \frac{1}{\beta_c(\beta_c + \beta_s)}\right] = \frac{2(2\beta_c + \beta_s)}{\beta_c(\beta_c + \beta_s)^2}, \qquad \mathrm{Var}(x^\star(0)) = \frac{\sigma^2 \beta_c (2\beta_c + \beta_s)}{2(\beta_c + \beta_s)^2}. \tag{C.1}$$

Combine. With $\mathrm{Var}(\theta(0)) = \sigma^2$,

$$\mathrm{Corr}^2 = \frac{\big(\sigma^2 \beta_c/(\beta_c + \beta_s)\big)^2}{\sigma^2 \cdot \sigma^2 \beta_c (2\beta_c + \beta_s)/\big(2(\beta_c + \beta_s)^2\big)} = \frac{2\beta_c}{2\beta_c + \beta_s}, \qquad \eta_S(0) = \frac{\beta_s}{2\beta_c + \beta_s} = \frac{\ell_c}{\ell_c + 2\ell_s}.$$

□

Monotonicity in the decision-relevance length. For the positive-radius regime $vT > 2\ell_s$, Theorem 3 can be written as

$$r^\star = \frac{1}{2} \int_{2\ell_s}^{vT} \frac{\ell_c}{\ell_c + s}\, ds. \tag{C.2}$$

With vT and ℓ_s fixed, the integration limits do not change and $\ell_c/(\ell_c + s) = 1 - s/(\ell_c + s)$ is strictly increasing in ℓ_c for every $s > 0$. Thus $r^\star$ increases strictly with ℓ_c in this regime, proving the fourth comparative static.

D Numerical reproducibility details

This section gives the settings and estimators for the four figures in Section 5. The crossover uses the exact coupled fresh-local rule; the radius curves use scalar variance composition; and the ring and four-geometry comparison use agentwise Gaussian conditioning with $Q = I$. Normalized scalar and graph losses use analytical open-loop regret R_∞. The crossover instead uses paired sample means. Pointwise Monte Carlo means have 95% normal intervals $\bar{L} \pm 1.96 s_L/\sqrt{N}$, with sample standard deviation s_L. Conditioning uses linear solves rather than explicit matrix inverses.

Table 1: Settings for the four numerical illustrations. The shared-resource model fixes $q = 1$; graph covariances have average marginal variance one.

Figure	principal sweep	fixed setting/size	samples
1	9 γ/q values in $[0, 10]$	$n = 5, 10, 25, 100$	30,000/setting
2	$r/\ell_c \in [0, 2]$ (1,001 points)	3 regimes; $D/\ell_c = 2$	50,000/point
3	$\ell_c = 0.7, 2, 6$	64-node ring; $T = 6$	12,000/check
4	all integer radii; T, ℓ_c	four 64-node graphs	exact

Fresh-local/stale-global crossover (Fig. 1). We fix $q = 1$, use $n \in \{5, 10, 25, 100\}$, and sample independent standard-normal state components ($\sigma^2 = 1$). The nine coupling ratios are $\gamma/q \in \{0, 0.05, 0.1, 0.2, 0.5, 1, 2, 5, 10\}$. At each setting, 30,000 independent current-state draws give paired open-loop and fresh-local losses, using $u_i = \theta_i/(q + \gamma)$. Writing their means as $o = \widehat{R}_\infty$ and $l = \widehat{R}_{\mathrm{loc}}$, the estimator is $\widehat{\rho}^\star = \frac{1}{2}\log[o/(o-l)]$. Its 95% delta-method interval is $\widehat{\rho}^\star \pm 1.96\sqrt{g^\top C g/N}$, where C is the sample covariance of the paired losses and

$$g = \left(-\frac{l}{2o(o-l)}, \frac{1}{2(o-l)} \right)^\top .$$

The paired covariance is retained; $o - l \le 0$ is rejected rather than clipped. The zero-coupling boundary is zero and is omitted from the logarithmic plot.

Three radius regimes (Fig. 2). The scalar calculation uses $\eta_S(r) = \eta_0 e^{-2r/\ell_c}$, $\tau(r) = r/v$, $\ell_c = 1$, and $D = 2$, with $(\eta_0, vT) = (0.25, 0.2), (0.7, 4), (0.9, 100)$. The radius grid has spacing 0.002. At each point, 50,000 pairs of independent standard normals are scaled by the square roots of the normalized delay and limited-scope terms in (28), then added and squared. This checks scalar variance composition; it does not simulate a spatial field.

Decision relevance on a ring (Fig. 3). The graph and covariance are fixed on a 64-node ring, with $T = 6$ and latency 0.12 per hop. For $\ell_c \in \{0.7, 2, 6\}$ the operator is $K_{ij} = e^{-d_G(i,j)/\ell_c}/\|e^{-d_G(i,\cdot)/\ell_c}\|_2$. With $Q = I$, exact Gaussian conditioning gives spatial omission; 12,000-sample checks evaluate selected radii. Each operator has its own analytical $R_\infty = \frac{1}{2}\operatorname{tr}(K\Sigma_S K^\top)$. Raw state correlation is computed separately from the fixed covariance.

The covariance in Eq. (39) uses $(\kappa_s, \nu) = (0.75, 1.5)$ and is scaled to average marginal variance one. Every integer radius from 0 through the ring diameter 32 is evaluated. For each operator, Monte Carlo checks use radii $\{0, r^\star, \min(r^\star + 3, 32), 32\}$. Independent current-state Gaussian draws are passed through the exact conditional policies; summed squared prediction errors are divided by $\operatorname{tr}(K\Sigma_S K^\top)$ to estimate the omission share.

Finite noncanonical graphs (Fig. 4). The original Experiment 6 uses a 64-node path and cycle, an 8×8 grid without wraparound, and one connected random geometric graph. For the latter, NumPy's `default_rng(20260839)` draws 64 independent uniform positions in $[0,1]^2$. Edges join pairs within Euclidean distance 0.22; the construction increases this threshold by 0.02 until connected (0.22 suffices for this seed). Observation radii use hop distance in every graph. Their diameters are 63, 32, 14, and 8, respectively. There is no graph-ensemble average or Monte Carlo estimation in this experiment.

The covariance is proportional to $(0.8^2 I + L_G)^{-1.5}$, scaled to average marginal variance one. The target rows are $K_{ij} = e^{-d_G(i,j)/\ell_c} / \|e^{-d_G(i,\cdot)/\ell_c}\|_2$ and $Q = I$. For neighborhood $O_i = \mathcal{N}_r(i)$, the residual variance is $K_i \Sigma_S K_i^\top - c_i^\top \Sigma_{S,O_iO_i}^{-1} c_i$, where $c_i = \Sigma_{S,O_i,:} K_i^\top$; linear solves evaluate this expression. The sum divided by $\mathrm{tr}(K \Sigma_S K^\top)$ gives $\eta_S(r)$. We evaluate $1 - e^{-2ar/T}(1 - \eta_S(r))$ at every integer radius from zero through the diameter, using reference $(T, \ell_c, a) = (5, 2, 0.12)$.

The plotted sweeps vary $T \in \{1.5, 2.5, 4, 6, 9, 14\}$ or $\ell_c \in \{0.7, 1.2, 2, 3.5, 6\}$ alone. The original archive also retains per-hop delays $a \in \{0.05, 0.08, 0.12, 0.18, 0.28\}$ and covariance parameters $\kappa_s \in \{1.5, 1, 0.7, 0.45, 0.3\}$, each varied alone. Configuration `numerics/config/experiment_06.yaml` records experiment seed 20260828 and graph seed 20260839. The original driver is `numerics/experiments/exp06_general_graphs.py`, which calls `exp06()` in `numerics/experiments/run_all.py`. From the repository root, run `python scripts/verify_graph_restoration.py` to rerun that function in scratch space and compare both processed CSVs and all four raw graph archives under `numerics/results/`. The audit checks positive-definite covariances, conditioning through an independent residual-error calculation, and exhaustive integer-radius minima. Its report is `submissions/arxiv-v2/graph-restoration-verification.json`. The restored figure reuses the original PDF bytes.

Reproduction of Figures 1–3. The public TINA repository retains the original experiment identifiers. Figures 1, 2, and 3 correspond to `numerics/config/experiment_02.yaml`, `numerics/config/experiment_03.yaml`, and `numerics/config/experiment_05.yaml`, with seeds 20260824, 20260825, and 20260827, respectively. Their drivers are `numerics/experiments/exp02_local_global_crossover.py`, `numerics/experiments/exp03_radius_regret.py`, and `numerics/experiments/exp05_decision_relevance.py`. Run metadata, processed results, and selected raw arrays are stored under `numerics/results/`; not every random draw is archived. `docs/numerical_method.md` and `docs/experiment_parameters.md` specify conditioning safeguards and environment versions; `numerics/README.md` gives reproduction commands.